


Stage-3 centreline comparison: degraded vs. clean input masks



Same 125 scenes in both panels (25 per density level), paired scene-by-scene. Markers show the per-density mean common-fragment F_1 (individual scenes not plotted); error bars are 95% percentile bootstrap CIs over scenes (seed 20260729, 10000 resamples). (a) mask_w1.png -- the primary, predefined input (degraded/UNet-realistic; mean foreground width ~ 1.8 px, area/skeleton-length). (b) mask_clean.png -- a secondary, non-predefined analysis on the undegraded, exactly-1px axis mask.