

PLECTA: Overlap-aware instance reconstruction of dense filament networks from binary masks

Oday Allan^a Stefan Luding^a Igor Ostanin^{a,*}

^a Multi-Scale Mechanics (MSM), Faculty of Engineering Technology, University of Twente, Enschede, The Netherlands

[To do: Add ORCIDs.]

Abstract

Dense filament networks resist instance segmentation because their two failure modes pull in opposite directions: crossings join distinct objects into one connected component, while gaps in the segmentation divide a single object into several. We present PLECTA, a deterministic method that converts a binary filament-axis mask into overlap-aware two-dimensional instances. It decomposes the skeletonised mask into arms and junctions, alternates exact per-junction continuation matching with globally gated gap matching, and re-estimates tangent, chord, and curvature evidence from the assembled chains, so that a crossing pixel may belong to more than one output instance.

Evaluation used a method-independent partition of each input skeleton into common fragments. On 128 independently seeded held-out scenes from the same procedural generator, PLECTA achieved mean pairwise $F_1=0.854$ (density-stratified 95% bootstrap CI [0.845, 0.864]), precision 0.868, recall 0.843, recovery 0.799, and adjusted Rand index 0.851, with no failed scenes. Mean F_1 decreased from 0.911 at 20% coverage to 0.775 at 60%, identifying dense crossings as the main remaining regime of error. Against a greedy minimum-turning-angle continuation rule on identical masks, PLECTA gained +0.145 F_1 [+0.129, +0.162] over 50 scenes, and +0.478 over a released implementation run from its own source at a development-selected configuration. Development-only ablations showed that chord-turn evidence, gap bridging, and junction-topology consolidation made the largest contributions.

CNT-sheet scanning electron microscopy is the proof-of-concept application: an upstream synthetic-data-driven CycleGAN/U-Net route prepares the binary masks that let the method run on real micrographs, while SEM width measurement and ribbon rendering are downstream stages that leave the grouping fixed. The evidence supports same-generator synthetic generalisation and a two-field demonstration on real SEM data, not general real-image or three-dimensional reconstruction claims.

Keywords: filament networks; instance segmentation; graph matching; geometric reconstruction; carbon nanotubes; scanning electron microscopy.

[Igor 1] I must say I'm not quite happy with FilaSeg - we need to come up with a better name of the software. I think we do more than just "Seg", I'd avoid that in the name anyways. I would not go towards something trivial like FIBER or STRING either - not enough distinctive. I find attractive something like REUSS, OCCAM or COSSERAT - with reference to physics/mechanics besides the meaning in abbreviation. But that's your baby of course, your word will be the last.

[Response] Addressed. The method is now PLECTA, after Latin *plecta*, a braid or twisted rope. It carries a meaning tied to what the method does, avoids *Seg*, and is not generic. No expansion is forced onto the letters.

*Corresponding author: i.ostanin@utwente.nl

1 Introduction

[Igor 2] The paper definitely should have one (and only one) focused, bright central idea. My preference - we present a black box that magically transforms microphotographs of entangled bundles of sticky elastic rods into high-resolution, segmented 3D models. We can frame it as the story about CNTs and make it carbon-specific or talk about arbitrary systems (spagetti, straw, bone tissue etc) and make it an image analysis paper. But we certainly should not pick separate elements of the pipeline and discuss only those - this way the value of the paper becomes incrementally small.

[Response] Open decision, not yet settled. The single central idea here is *recovering individual filament identity through crossings*, and the paper is built around it end to end: micrograph, mask, instances, geometry. Two constraints pull against the wider framing. The reconstruction is projected and two-dimensional, so *3D models* would overclaim unless we add depth evidence; and the upstream segmenter is a separate contribution with its own evaluation, so folding it in would mean defending two methods at once. The current draft therefore keeps the CNT story as the application and the identity problem as the idea. Please say if you would rather we widen it.

Separating the individual objects in an image of a dense filament network is an instance-segmentation problem with an awkward geometry. The objects are long, thin and open, they touch and cross and occlude one another, and what distinguishes them is local continuation rather than shape or extent. This paper presents a method for that problem. Its proof-of-concept application is scanning electron microscopy (SEM) of carbon-nanotube (CNT) sheets, whose mechanical and electrical response depends on the geometry of the bundle network: length, width, orientation, curvature and contact topology [1–5]. SEM resolves that projected network, but routine image analysis often stops at coverage or orientation distributions [6, 7], whereas instance-resolved analysis requires each visible bundle to remain addressable across gaps, touches and crossings.

That requirement is distinct from semantic segmentation. U-Net and its self-configuring variants can identify filament foreground [8, 9], but a binary foreground mask does not encode object identity: a faded segment splits one filament into several connected components, and a crossing joins different filaments into one. The ambiguity is greatest in dense networks, where several local continuations are geometrically plausible. Compact-object instance priors (bounded blobs, centre-directed flows, star-convex shapes) are poorly matched to long open curves that may overlap [10–12]. Ridge and vesselness filters enhance curvilinear signal but do not assign identity across junctions [13–15], and fibre extraction packages such as FIRE, CT-FIRE, SOAX, FiNTA and DiameterJ report network statistics rather than resolved instances [16–20].

Recovering individual filaments through crossings is itself established prior art, and the design space is well explored. Basu et al. build a minimum spanning tree over centreline particles, delete every bifurcation node, and recombine the segments with a binary integer program over chains of at most three, priced by the integrated squared curvature of a fitted spline; a segment may stay unlinked, and the program is re-solved after refitting until the segment set stops changing [21]. DeFiNe consumes a pre-extracted weighted graph and solves a filament cover as a linear program minimising path roughness, optionally letting several paths cover one edge, so overlaps are representable [22]. SIFNE excises each crossing with its neighbourhood, estimates tip directions towards the local centre of mass, and pairs termini greedily in priority order inside a fan-shaped region under angular and continuity gates, leftover tips becoming ends [23]. De et al. resolve crossovers by partitioning a directed graph [24], and Wegmayr et al. group skeleton branches by lowest mutual angle and learn pixel embeddings for tangled microplastic fibres, sharing junction pixels in principle but assigning them arbitrarily [25]. DNAi detects junctions morphologically, clusters them by single linkage, erases a disk the width of the fibre, and enumerates the pairings at junctions of degree three or four from endpoint angles traced over a short span [26]. GraFT builds a graph from a Frangi-filtered skeleton, derives an angle threshold per component, and

traces filaments by constrained depth-first search and a greedy longest-first path cover [27]. Learned alternatives include orientation-aware terminus pairing [28], FibeR-CNN, and recurrent pick-and-trace models [29, 30].

Skeleton graphs, junction excision, continuation angles, distance gates, combinatorial optimisation, iterative refitting and overlapping filament output are therefore all prior art, and none is claimed here. Two distinctions are narrower and survive checking against the sources. First, the methods above either pair endpoints greedily in priority order [23, 25, 27] or enumerate the pairings of a junction only up to degree four [26]; PLECTA instead solves each junction of any degree as one maximum-weight matching in which leaving an arm unmatched is an explicit option at a fixed price, rather than a parity leftover. Second, the methods that optimise globally require intensity evidence or a pre-extracted weighted graph [21, 22], whereas PLECTA reads a binary mask alone. The closest single prior method is that of Liu et al., which likewise works from a mask, reconnects across gaps and permits overlapping output, but pairs termini greedily and depends on a trained orientation network [28].

We introduce PLECTA, a deterministic graph-matching method for dense filament networks, named after the Latin *plecta*, a braid. PLECTA decomposes a skeleton into arms, endpoint stubs, and merged junction nodes. It then alternates exact per-node continuation matching with globally gated gap matching while tangent and curvature estimates expand from local arms to assembled chains. Its output is a set of projected 2-D instance layers in which crossing pixels may be shared.

Reaching real micrographs takes one more component. A complete CNT workflow combines semantic segmentation with algorithmic tracing and transverse intensity profiles [31], and we divide that workflow in two. Pipeline A turns a raw SEM micrograph into a binary filament-axis mask by a synthetic-data-driven CycleGAN/U-Net route [32, 33]; Pipeline B is PLECTA. Pipeline A is what makes the proof of concept on real data possible at all, since without it the method could only ever be exercised on masks drawn by hand or emitted by a generator. The division also lets grouping be measured on its own, against references whose identities are known exactly, and keeps width rendering from moving an identity score. Pipeline A carries its own evaluation [33] and is not re-evaluated here. We evaluate PLECTA on 128 unseen scenes from the development generator and on 50 paired clean and degraded scenes, and demonstrate the two pipelines together on two annotated CNT-sheet SEM fields.

2 Materials and methods

2.1 Scope and workflow

The formal input is a binary filament-axis mask $M \in \{0, 1\}^{H \times W}$, from which PLECTA reconstructs an overlap-aware collection $\{I_1, \dots, I_N\}$ with $I_n \subseteq M$ and a crossing pixel possibly belonging to several instances. The evaluated grouping stage uses only M ; a grayscale SEM image is optional downstream input for width and brightness measurement and cannot revise the fixed grouping.

Figure 1 places the method in the two-pipeline workflow of Section 1: Pipeline A produces the mask from a micrograph, and Pipeline B, which this paper evaluates, is PLECTA. Its five geometric steps read mask geometry alone, and the two matching steps are re-made over eight rounds against frames re-fitted along the chains assembled so far.

What makes the task hard is visible in a single crossing (Figure 2). Where two filaments meet, the mask holds one connected blob from which several arms leave, and nothing local states which arm continues into which. The output is a set of layers rather than a partition, so a pixel inside that blob is counted in both instances that claim it.




Figure 1. Where PLECTA sits. Pipeline A is the upstream binary segmentation of SEM micrographs; Pipeline B is PLECTA, and the held-out result measures the dashed green box only. The mask entering Pipeline B came from Pipeline A for the SEM fields and from the procedural generator in every synthetic experiment. Colours: blue, a learned component; green, a PLECTA geometric step; orange, a dependency on the grayscale image; dark, a data artefact.

2.2 Synthetic data and evaluation split

Procedural scenes preserve the ordered centreline $\mathcal{C}_k$ and rendered support Ω_k of every filament k ; supports remain separate, so several filaments may own the same crossing pixel. Besides the exact overlap-aware reference, the generator provides an undegraded one-pixel axis mask, a degraded thin mask containing small gaps and irregularities, a combined thick foreground, and an SEM-like image. The degraded one-pixel axis is the input evaluated throughout (Figure 2b,c). Areal coverage is the union area of the thick supports divided by image area, not the occupancy of the thin input mask, and Section 3.2 shows it is a proxy for difficulty rather than a cause of it.

Development used 84 scenes spanning 20–60% coverage. The held-out set comprised 128 different seeds, 32 scenes at each of 20, 30, 40, and 60% coverage. It is an independent sample from the same generator, not an independent distribution. A separate 125-scene locked set remains unevaluated. The compact seed manifest records every seed block and an executable verifier checks counts, disjointness, and available scene metadata.

A separate controlled robustness set of 50 scenes, ten at each of 20, 30, 40, 50, and 60% coverage, was scored once from each geometry’s clean axis and once from its degraded mask, after the configuration was fixed.




Figure 2. The task, and what the output is. (a) A four-arm crossing, junction cluster in red. (b) The evaluated input: one degraded 1 px binary axis mask. (c) The overlap-aware reference for that scene, 40 instances. (d) The crossing after PLECTA; split squares mark shared pixels. (e, f) The two instances separately, both holding the junction cluster. (g) Within this 52×52 crop each instance has 54 pixels, 7 of them shared. Instance colours are arbitrary and carry identity only within a panel.

2.3 CNT SEM data and mask sources

Two manually annotated CNT-sheet SEM fields were available for exploratory analysis: B58_100 (1024×1536 px, 280 filaments) and B58_110 (1024×1536 px, 430 filaments). The nominal horizontal field width is $1.66 \mu\text{m}$; project metadata associate it with 1536 px, or approximately 1.08 nm px^{-1} . The annotation is one human interpretation of projected, locally ambiguous crossings, and both fields informed development.

Two inputs were derived per field. The first was the skeleton of the manual instance annotation. The second came from Pipeline A: an nnU-Net trained on paired synthetic images and masks, with CycleGAN translation as the route from procedural masks to SEM-like training images [32, 33]. Together the two conditions separate what grouping can do from what the mask allows it to do. B58_100 appeared in that model’s training data, whereas B58_110 was verified absent from the relevant training set, so only B58_110 provides a training-clean observation of mask transfer, and $n = 2$ precludes inference.

[Needed before submission: A fuller account of Pipeline A belongs here, to be written by the author: the CycleGAN configuration used to translate procedural masks into SEM-like images, the paired synthetic corpus it produced, the nnU-Net variant and training split trained on that corpus, the optimisation settings, and how the predicted masks were selected and post-processed before entering PLECTA. Take these from the model archive rather than reconstructing them.]

Mask quality was compared after skeletonisation with a symmetric 3-px tolerance, which avoids rewarding mask thickness, and reconstruction scores were computed within each mask’s own visible common-fragment population.

[Needed before submission: Author metadata: SEM instrument and detector,

accelerating voltage, working distance, specimen preparation, CNT material identity, annotation protocol, and confirmation of the 1536-pixel calibration against the stored image dimensions.]

2.4 PLECTA reconstruction

2.4.1 Skeleton graph and local frames

PLECTA binarises M at > 0 , skeletonises it, and iteratively removes spurs no longer than 3 px. Skeleton pixels with degree at least three form junction clusters, and adjacent clusters are merged into nodes. Connectors up to 5 px are absorbed into one node; connectors up to 14 px merge nearby nodes but remain real arms; free debris up to 2 px is absorbed. Each remaining connected path is stored as an ordered arm with one stub at each end.

A stub frame contains its tip, outward tangent, and signed curvature, estimated by an arclength polynomial fit and considered reliable only with at least three samples spanning at least 3 px. Local frames use 24 px of support; once links form chains, the support expands to 55 px, so that a stub too short to have a direction of its own borrows one from the chain it has joined.

Link cost combines three angular terms, curvature continuity and, for gaps, a length penalty (Figure 3). The first angular term is the direct reversal angle between the two outward tangents; it ignores the chord between the tips, which is what keeps it meaningful when those tips are three pixels apart and the chord direction is quantisation noise. The other two are the turns from each tangent onto that chord, measured at their own vertices; unlike the reversal angle these see lateral offset, so two parallel arms that never meet are penalised. Because it is the chord terms that degrade over short separations, a fade factor withdraws them as the tips approach rather than letting an unstable angle dominate the cost.




Figure 3. The stub frame and the terms of the pairing cost. (a) The least-squares arclength fit over window w , read outward from the tip at $s = 0$, giving outward tangent $\mathbf{t}$ and signed curvature κ ; the fit is quadratic when the span is at least 18 px and linear otherwise. (b) $\theta = \arccos(\mathbf{t}_a \cdot -\mathbf{t}_b)$, drawn where the reversed tangent lies and again from the common origin where the angle is read. (c) The chord turns φ_a, φ_b onto the chord $\mathbf{u}$ joining the tips. (d) The fade $q(d)$ against tip separation; triangles mark the six real separations at the crossing of Figure 5.

2.4.2 Candidate gates and exact matching

Each of eight rounds first resolves junctions and then gaps. At a junction, candidates must share the node, use different arms, and have finite frames and cost. Exact maximum-weight partial matching then selects disjoint pairs while allowing stubs to remain unmatched. Leaving a stub unmatched is priced rather than forbidden, and that price also fixes admissibility: at full strictness a junction edge is offered only when its cost falls below 1.24, twice the unmatched price. Gap candidates are constructed only from still-unmatched, reliable stubs. They must belong to different arms and different junctions, lie within 85 px, have direct tangent mismatch at most 0.40 rad, and have chord-turn mismatch at each end at most 0.28 rad; global exact matching then accepts only edges costing below 0.60. The unmatched prices and angular limits are annealed linearly from 85% to full strictness over the eight rounds, and the eighth, fully strict round is the one returned.

Figure 4 enumerates the outcomes this admits. Every stub is either paired or pays the unmatched price, and there is no third possibility. The priced option is what lets a filament genuinely end, and what stops a T-junction from inventing a continuation for the arm that has no partner.




Figure 4. Every decision the linker can make. (a) A junction pairing accepted below the admissibility limit $2p_i$. (b) A competing pairing at the same node refused, its cost exceeding $2p_i = 1.24$. (c) A stub left unmatched because paying $p_i = 0.62$ beats any edge it has. (d) A gap joined between two stubs the junction stage left free, under four gates at once. Idealised drawing; measured values are in Figures 3 and 5.

Solving a junction as one matching rather than one edge at a time is the distinction claimed in Section 1, and Figure 5 shows what it buys on real nodes. A configuration is scored as the sum of its link costs plus the unmatched price for every stub left free, and a maximum-weight matching with edge weight $p_i + p_j - C_{ij}$ minimises that total exactly. A locally attractive edge can therefore be declined, as at the three-arm node of Figure 5e. Joint and sequential solutions coincide at plain crossings and diverge only as nodes grow dense: over 20 development scenes they never disagreed at nodes of degree 2–4 ($n = 542$), disagreed at 22 % of nodes with 5–8 stubs ($n = 187$) and at 69 % of nodes with 9 or more ($n = 233$), where the joint solution was better by a median of 0.90 in total cost.

2.4.3 Instance output

Connected components of accepted arm links define instances. PLECTA paints every arm pixel and every junction cluster touched by a chain into that instance, so intersecting instances share crossing pixels. An isolated one-arm component shorter than 6 px is omitted. Accepted




Figure 5. The pairing at a crossing, solved jointly and exactly; every number shown is measured. (a) A four-arm crossing (development scene `cov20/synth_0001`, node 9), arms named *a* to *d* anticlockwise by outward tangent, junction pixels red. (b) All six pairwise costs; open bars are the four refused by $C_{ab} < 2p_i = 1.24$. (c) Every admissible configuration, scored jointly; the boxed one is returned. (d) The result: two instances sharing the seven junction pixels. (e) A three-arm node (`cov40/synth_0001`) at which an admissible edge costing 1.01 is declined: taking it scores 1.63 overall, against 0.77 for pairing the other two arms and paying to leave the third free. (f) Joint against sequential cheapest-edge-first solutions, by node degree, over 20 development scenes.

gap links determine identity but are not painted into the fixed centreline output. All effective parameters, including values previously inherited as code defaults, are materialised in the fixed JSON configuration.

2.5 Optional SEM characterisation and ribbon rendering

Two optional stages act after grouping, both off by default and both governed by limits recorded in the fixed configuration. Perpendicular SEM profiles, sampled along each ordered chain away from junctions, attach apparent width, brightness, uncertainty, and a quality flag to an existing instance without altering its arm membership. Ribbon rendering then interpolates accepted gaps, resamples and smooths the centreline, and draws the chain at its fitted width, excluding junction blobs. Silhouette absorption, which would union rendered instances, has threshold zero by default; neither it nor any post-hoc continuation merge contributed to the held-out result, and neither is part of the evaluated algorithm.

2.6 Metrics and statistical analysis

The primary endpoint is common-fragment pairwise F_1 . For each scene the visible input skeleton is partitioned once into method-independent elementary fragments, and each fragment receives a procedural reference identity and a PLECTA instance identity. A pair of fragments is a true positive when it shares both identities, a false positive when it joins different reference filaments, and a false negative when it separates fragments of one. Thus

$$P = \frac{TP}{TP + FP}, \quad R = \frac{TP}{TP + FN}, \quad F_1 = \frac{2PR}{P + R}. \quad (1)$$

This is a clustering metric related to the Rand family [34, 35]. It scores grouping rather than

Table 1. Held-out reconstruction scores overall and by target areal coverage. VI denotes variation of information; lower VI is better, whereas higher values are better for the other metrics.

Group	n	F_1	Precision	Recall	Recovery	ARI	VI split	VI merge
All	128	0.854	0.868	0.843	0.799	0.851	0.256	0.216
20%	32	0.911	0.928	0.896	0.918	0.907	0.137	0.093
30%	32	0.874	0.892	0.858	0.828	0.870	0.211	0.161
40%	32	0.858	0.864	0.853	0.796	0.855	0.242	0.224
60%	32	0.775	0.787	0.765	0.656	0.772	0.436	0.388

rendered width, and weights filaments with many fragments more strongly because their pair count is quadratic, so we also report adjusted Rand index [36], variation of information split and merge components [37], and the fraction of reference filaments recovered by one dominant predicted instance at 0.8 purity and coverage.

Held-out means and density strata are accompanied by 95% nonparametric bootstrap intervals (20,000 replicates; seed 20,260,810), resampled with stratification by density; the clean-minus-degraded analysis resamples paired scene differences within density. Active-component ablations are one-at-a-time development experiments on all 84 development scenes and do not reuse the held-out set.

3 Results

3.1 Held-out synthetic reconstruction

PLECTA completed all 128 held-out scenes with 0 failures. Mean common-fragment pairwise F_1 was 0.854 (density-stratified 95% bootstrap CI [0.845, 0.864]); mean precision, recall, reference-instance recovery, and adjusted Rand index were 0.868, 0.843, 0.799, and 0.851, respectively (Table 1). These scenes used unseen seeds from the same procedural generator as development data, so this is same-generator generalisation and not distribution transfer (Section 4.2).

Performance decreased with scene density: mean F_1 fell from 0.911 [0.888, 0.933] at 20% coverage to 0.775 [0.760, 0.790] at 60%, the 30% and 40% means being 0.874 and 0.858. The same decline appears on the paired robustness set, where each geometry is scored from both input conditions (Figure 6): the two curves separate at low coverage and meet at 50–60%, degradation costing least exactly where the score is lowest, because there the tangles rather than the gaps set the difficulty. Per-scene runtime had mean 2.54 s, median 1.12 s, 90th percentile 6.42 s, and maximum 34.06 s; these timings are hardware-specific.

3.2 Which density axis matters

Coverage is the fraction of the frame filled by the rendered bundles, so it rises both when filaments are added and when the same filaments are drawn thicker, and only the first changes the grouping problem. A factorial separates the two. Twelve geometries, three centreline length densities crossed with four seeds, were each rendered at bundle widths of 6, 11 and 16 px. Within one geometry the input axis mask and the overlap-aware centreline reference are byte-identical across the three widths, verified by SHA-256, while the union of the thick supports changes: areal coverage varies by up to a factor of 2.49. Any movement in the score across those three renderings would mean that the method or the metric reads the silhouette.

Across all 12 geometries the within-geometry F_1 range was 0.0000, that is, exactly zero in 12 of 12 (Table 2). Areal coverage therefore cannot causally affect the grouping score; PLECTA reads the axis mask, and the metric scores centreline fragments, so neither sees the width the bundles




Figure 6. Coverage response on the 50 paired scenes, ten per coverage level, each scored twice: from the degraded 1 px axis mask that the paper evaluates, and from the generator’s undegraded axis. Lines are means of per-scene common-fragment F_1 ; bands are 95 % percentile bootstrap intervals for the mean, resampled within coverage stratum (20 000 replicates, seed 20260811). Target areal coverage is the generator setting used to dial in centreline density.

Table 2. Bundle width against centreline length density, 36 development scenes. Within each row the three bundle widths share one byte-identical axis mask and one reference, so the spread in areal coverage carries no score difference. Length density is centreline pixels per image pixel.

Centreline length density	n	Areal coverage	Bundle width	F_1
0.0193	12	0.118–0.309	6, 11, 16 px	0.920
0.0372	12	0.218–0.505	6, 11, 16 px	0.878
0.0548	12	0.305–0.644	6, 11, 16 px	0.836

were drawn at. Centreline length density does move it: mean F_1 fell from 0.920 to 0.836 as centreline coverage rose from 0.0193 to 0.0548. The coverage bands of the three length densities also overlap, an areal coverage near 0.31 occurring at all three, so coverage does not identify a difficulty level on its own.

Density strata reported by target areal coverage should therefore be read as a proxy for the length density that accompanies them under this generator, and a mask drawn at a different thickness, as an upstream segmenter may well produce, does not by itself move the grouping score.

3.3 Comparison with greedy continuation

The contribution claimed in Section 1 is that a junction is better solved as one exact matching, with an explicit priced option to leave an arm unmatched, than by pairing endpoints greedily in priority order, the rule used by SIFNE, by GraFT’s path cover, and by the continuity baseline of Wegmayr et al. [23, 25, 27]. The comparator therefore skeletonises the same mask, prunes spurs, splits the skeleton into branches, and greedily accepts the lowest-turning-angle pairing between branch ends within a distance and a turning-angle gate. Both gates were selected on the same 84 development scenes PLECTA was configured on, by sweeping a grid and taking the best mean F_1 , so neither method is advantaged by tuning effort. Both received identical masks and one scorer, on 50 scenes generated after both configurations were fixed.

Table 3. PLECTA against a greedy minimum-turning-angle continuation baseline on 50 scenes with identical input masks and one scorer. Differences are paired, with density-stratified 95% bootstrap intervals.

Measure	PLECTA	Greedy continuation	Difference
Common-fragment F_1	0.822	0.676	+0.145 [+0.129, +0.162]
Precision	0.831	0.785	+0.046 [+0.029, +0.062]
Recall	0.814	0.599	+0.215 [+0.196, +0.235]
Reference-instance recovery	0.765	0.542	+0.223 [+0.201, +0.245]
Adjusted Rand index	0.817	0.670	+0.148 [+0.131, +0.165]
Variation of information (total) ↓	0.573	1.036	-0.463 [-0.509, -0.418]

PLECTA scored 0.822 against 0.676, a paired difference of +0.145 (95% CI [+0.129, +0.162]), and was better in 49 of 50 scenes (Table 3). Adjusted Rand index and variation of information rank the two the same way, so the result does not depend on the in-house measure. The gap is widest in recall and in reference-instance recovery, the expected signature: a greedy pass commits to the locally straightest continuation and cannot revise it when that choice strands a better partner, so single filaments are left divided. On clean axes the difference was +0.128 [+0.113, +0.144], so it is not an artefact of mask degradation.

3.4 Comparison with a released implementation

The greedy rule above is a reimplementaion, so we also ran a published method from its own source. DNAi [26] post-processes a binary segmentation by the sequence described in Section 1. Its segmentation network was bypassed and only that stage was used, so grouping is not confounded with segmentation quality, and both methods received the same masks and the same scorer.

At its released defaults DNAi scored $F_1=0.195$ on the degraded masks, which is not a fair measure of it. The cause is its single-linkage clustering of junction pixels at 10 px: our crossings are closer together than that, so the clustering chains, and at 60% coverage the largest cluster spans 89 px and is replaced by one small erasure disk, leaving the crossings along it unseparated and the whole tangle a single instance. Reducing the clustering distance to 2 px on development scenes raised development F_1 from 0.211 to 0.378 and cut variation-of-information merge from 2.07 to 0.66, confirming the mechanism. Dilating the input to a thicker mask was also tested across 41 development configurations and made DNAi worse at every setting, so it was not used.

At its best development-selected configuration DNAi scored 0.343 against PLECTA’s 0.822, a paired difference of +0.478 (95% CI [+0.456, +0.501]) over 50 scenes; on clean axes the difference was +0.414 [+0.392, +0.436]. The gap widens monotonically with density (Figure 7). Both methods must be read against the one-instance-per-connected-component floor at their own density, since that floor is itself strongly density-dependent and a pooled value would flatter any method at 60% coverage.

This is a limitation of domain and density, not a general ranking. DNAi is built for DNA-fibre micrographs, which are sparse: its own bundled test image contains five junction clusters, whereas our scenes contain between 23 and 183, and where the network is sparse and the axis clean it recovers 0.660 at 20% coverage. Two structural differences explain the rest. It commits to a pairing at every junction of degree two to four with no cost threshold, so it cannot decline a continuation the way a priced unmatched option can; and it resolves only junctions of degree at most four. That cap is easy to underrate by counting junctions, since junctions of degree five or more are only about a tenth of the clusters here. Weighted by the evidence the metric actually counts, they carry 37.2 % of same-filament fragment pairs on the degraded masks and 55.2 % on




Figure 7. PLECTA, the greedy continuation rule and DNAi across coverage, on the same 50 scenes with identical masks and one scorer, for the degraded and clean axis conditions. Plain means, ten scenes per point, no interval and no inference. The connected-component floor is drawn per density and falls from 0.146 to 0.041 across this range. DNAi is shown at its released defaults and at the best configuration selected on development scenes.

the clean ones, because a single such junction severs every pair that spans it.

That raises the question of how much of the margin is capability rather than judgement. Restricting the score to the evidence a degree-limited, gap-blind method can reach answers it. On clean axes the difference falls from +0.414 to +0.284 [+0.262, +0.307], still favouring PLECTA in 49 of 50 scenes and at every coverage level. DNAi gains far more from the restriction than PLECTA does, which is what shows the restriction genuinely removes its handicap. So roughly 31 % of the margin is coverage of cases the comparator cannot attempt and 69 % is the outcome of decisions both methods make. Consistent with that, after its clustering is repaired DNAi’s remaining errors are dominated by division rather than merging, and a third of those divisions occur at junctions of degree four, which it does attempt: its matcher commits to a pairing there with no rejection threshold. Section 4.2 states what both comparisons leave untested.

Individual scenes show what the score corresponds to. Figure 8 takes the scene whose F_1 is closest to the median at three coverage levels, so the rows are typical rather than selected, and shows the evaluated output beside the same instances after the optional rendering stage. The grouping is scored on the centrelines in the third column; the fourth adds the fitted widths and is what the full pipeline delivers. It is included because the rendered form is what a downstream user sees, and because the comparison with the reference is only legible once the instances are drawn at width. Errors survive rendering unchanged: it redraws the instances the grouping produced and cannot merge or divide them.

3.5 Input-mask robustness

On the separate 50-scene controlled set, mean F_1 was 0.822 from degraded axes and 0.844 from the corresponding clean axes. The paired clean-minus-degraded change was 0.0224 (95% CI [0.0113, 0.0333]). Modelled gaps and raster irregularities therefore caused a small resolved loss under this generator.

Mask thickness costs more than degradation does (Figure 9), which matters because it is the upstream segmenter, not the method, that fixes that thickness. Each width variant is scored against common fragments rebuilt from its own input mask, so the question asked is how well




Figure 8. Reconstruction at three coverage levels, each the scene closest to the median F_1 of its stratum. Columns: the degraded input mask, the overlap-aware reference, the evaluated centreline instances, and the same instances after optional width rendering. Only the third column is scored. Instance colours are arbitrary, carry identity only within a panel, and are not matched between columns.

the method groups the skeleton it is handed rather than how many pixels it was handed; the decline with width is a thicker mask skeletonising into a messier graph, not a change in the grouping problem.

3.6 Development ablations and downstream rendering

One-at-a-time ablations were evaluated only on the 84 development scenes (Table 4). Relative to the evaluated configuration, removing the junction chord-turn term changed F_1 by -0.1858, removing gap matching by -0.1394, and disabling connector-based junction-cluster consolidation by -0.0634. Every other component, chain-extended frames and the round schedule included, mattered less, the largest such change being -0.0221. Chord-aware continuation, gap matching, and junction-topology preconditioning are thus the largest tested contributors, but these are development values and not held-out effect estimates.

The core output is a grouping, not a set of connected curves: accepted gap links establish identity but are not painted, so on the development set 0.653 of core instances arrive as more than one disjoint pixel run and 0.663 contain a branch point. Any downstream step that assumes one connected curve per instance must account for this. Default Stage 4 ribbon rendering scored development $F_1=0.860$, a change of -0.0068, and reduced those fractions to 0.000 and 0.005, so rendering buys connectivity at a small cost in the identity score without changing chain membership. Enabling the separate 0.6 silhouette-absorption rule reduced F_1 further to 0.857, and because absorption can merge identities it remained disabled. None of these topology




Figure 9. How PLECTA responds to the input mask it is handed, on development scenes. (a) Nominal binary axis-mask width. (b) The degraded 1 px axis against the generator’s undegraded axis. Lines are means over the four development scenes at each coverage level; bands are the min–max range across those scenes, a description of spread and not an interval estimate. Exploratory, $n = 4$ scenes per point, no inference.

Table 4. One-at-a-time component ablations on the 84-scene development set. Negative ΔF_1 values denote losses relative to the evaluated configuration.

Development variant	F_1	ΔF_1	ARI
Fixed configuration	0.867	+0.000	0.864
No Gap Bridging	0.728	-0.139	0.722
Single Round	0.847	-0.020	0.844
No Chain-Extended Frames	0.846	-0.021	0.842
No Annealing	0.863	-0.004	0.860
No Curvature Term	0.863	-0.005	0.859
No Join_Px	0.845	-0.022	0.841
No Junction-Cluster Consolidation	0.804	-0.063	0.799
No Junction Chord-Turn	0.681	-0.186	0.675

measurements validates apparent width.

Two properties of the held-out result are invisible in a table of means (Figure 10). The first is the negative control of Section 3.2 seen scene by scene: within a fixed geometry the twelve lines are flat, so the coverage axis of Figures 6 and 7 stands in for centreline density rather than causing the difficulty itself. The second is that PLECTA’s errors are two-sided. A fragment F_1 can be reached one-sidedly, by a method that merges everything or one that divides everything, so it matters that PLECTA sits on the precision–recall diagonal. The connected-component floor does not: it recovers 0.850 of same-filament pairs at a precision of 0.047, because it merges every filament of a connected network into one instance. That placement is what makes it legible as a diagnostic rather than a competitor.

3.7 Two-field CNT SEM mask-source case

The two real SEM fields are the proof of concept for the whole route, and are exploratory development cases (Table 5). From the manual-derived axis, an oracle-input condition, PLECTA




Figure 10. Held-out spread, the width factorial, and the precision–recall plane. (a) All 128 held-out scenes, 32 per coverage stratum, stratum means printed above; bands describe spread across scenes and are not intervals for the mean. (b) Twelve development geometries, one line each; along a line the geometry is fixed and only bundle width changes, so realised areal coverage moves by up to a factor of 2.49 while F_1 moves by 0.0000. (c) The same scenes, one point each, against the one-instance-per-connected-component floor.

Table 5. Exploratory reconstruction from manual-derived and U-Net axes in two CNT SEM fields. Axis recall and precision compare the U-Net skeleton with the manual-derived skeleton using a symmetric 3-px tolerance.

Field	Input mask	Axis recall	Axis precision	F_1	ARI
B58_100	Manual-derived	–	–	0.882	0.882
B58_100	U-Net	0.782	0.753	0.438	0.435
B58_110	Manual-derived	–	–	0.837	0.836
B58_110	U-Net	0.729	0.887	0.452	0.449

reached $F_1=0.882$ on B58_100 and 0.837 on B58_110; with Pipeline A’s U-Net axes, whose recall and precision against the manual skeleton are given in the table, the values were 0.438 and 0.452. The distance between the two conditions is the practical dependence of the result on upstream mask quality. With two development fields, one of them training-contaminated for the U-Net condition, and a fragment universe that differs between the two mask sources, these values are descriptive and do not estimate real-image generalisation (Section 4.2).

4 Discussion

PLECTA addresses a narrow problem: assigning fragments of a binary filament-axis mask to overlapping projected instances. A held-out common-fragment F_1 of 0.854 shows that explicit geometry recovers substantial identity with no instance-labelled training data and no SEM intensity. What recovers it is the combination of local and accumulated evidence, neither of which suffices alone: a junction is resolved by mutual compatibility rather than by the cheapest available continuation, and the evidence for that judgement is re-read along the chains as they assemble.

The ablations locate that effect without proving each term independently necessary: the two large losses correspond to the two kinds of discontinuity, the junction and the missing mask segment, while accumulated context is worth distinctly less than either.

4.1 Density, masks, and optional rendering

Difficulty follows the density of centrelines rather than the amount of ink, so what costs the score is the number of locally plausible alternatives at a crossing, and the decline with coverage is a decline with the length density that accompanies it under this generator.

The rest of the sensitivity belongs to the upstream segmenter, which fixes the thickness of the mask and therefore how clean a skeleton PLECTA is given. PLECTA can bridge a short omission, but it cannot infer a filament absent from the mask, nor reliably reject a false axis introduced upstream. Supplying identical masks isolates grouping in an algorithmic comparison; it does not establish end-to-end SEM performance, which is why the two-field case is a demonstration of the joined pipelines rather than a measurement of them.

Ribbon rendering, being downstream of grouping, buys connectivity rather than identity: it delivers as single connected curves nearly all of the instances the core emits in several pieces, at a small cost in F_1 , because the core paints arms and junction clusters but not the gap links that carry identity across a break. Silhouette absorption differs in kind, since it changes the partition itself, and so stays optional. Width and brightness can be reported for characterised instances, but their physical accuracy requires separate references.

4.2 Evidence boundary

The principal limitation is distributional. The held-out scenes use independent seeds but the same procedural generator family used for development. Their confidence intervals quantify sampling variation within that family, not transfer to another simulator, specimen, microscope, or segmentation pipeline. The ablation and rendering studies are development-only and remain secondary to the fixed held-out result.

The real SEM study is deliberately a case study with two fields, not a population validation. Manual-derived masks are an oracle-like input condition: they test grouping when the axis evidence is curated, not automatic SEM-to-instance performance. B58_100 is training-contaminated for the U-Net condition and B58_110 documented as training-clean, so only the latter illustrates an unseen upstream mask, and one field cannot estimate generalisation. The two mask sources also generate different fragment universes, so their F_1 values describe performance under each input condition rather than comparing identical objects. More independent fields, sealed before configuration, and annotations from multiple readers are required.

Two further limits concern the output itself. PLECTA recovers projected two-dimensional identity: shared pixels encode overlap but not which filament passes above another, so no three-dimensional topology follows. And the grouping metric evaluates ownership of centreline fragments, which is not a validation of rendered width; quantitative diameter claims require matched physical measurements across the relevant scale, contrast, and specimen ranges. Pairwise fragment scores also emphasise instances containing many fragments, which is why object recovery and ARI are kept alongside them.

Both comparisons are bounded. The greedy rule is our reimplementing of a principle rather than of any particular paper, and DNAi was run outside the domain it was built for: its own bundled test image holds five junction clusters against 23 to 183 in ours, and part of its cost function reads a colour channel our masks do not have. What the two results support is narrow and mechanical, that pairing endpoints greedily, or committing to a pairing with no option to decline, loses identity in dense crossings. They do not support a claim that PLECTA is generally superior; that would need several released implementations run in their own regimes by their own authors.

An independent procedural benchmark is the practical next test. The `merged_lines` mask

generator in the CNT_SEM pipeline can export overlap-aware whole-bundle references and would add a different implementation and morphology model, but it should supplement rather than replace the current generator: its field-of-view preset is calibrated to a narrow coverage regime where the primary benchmark spans density. A multi-scene run should follow only once that calibration is fixed and its seeds separated from any configuration work, and no result from it is included here. Broader validation should then extend to independent real masks and other filament domains while keeping input fragments and scoring rules fixed.

5 Conclusions

PLECTA reconstructs overlapping filament instances from a binary axis mask by alternating exact junction matching with gated gap matching as local arm geometry expands into chain-level evidence. The fixed method achieved common-fragment $F_1 = 0.854$ on 128 independently seeded held-out scenes from the same procedural generator, with performance declining as network density increased. On identical masks it gained +0.145 F_1 over a greedy minimum-turning-angle continuation rule, the rule shared by several published linkers, and +0.478 over a released implementation run from its own source; both support solving each junction exactly rather than pairing endpoints in priority order, though the second is a limitation of domain and density rather than a general ranking. Development ablations identify junction chord-turn evidence, gap bridging, and junction-topology consolidation as the principal contributors, and optional image-informed rendering improves geometric presentation while remaining separate from the grouping claim.

The proof of concept on carbon-nanotube sheets joins the two pipelines: an upstream synthetic-data-driven segmenter prepares the mask, and PLECTA resolves it into instances. Two SEM fields demonstrate that route under manual-derived and U-Net masks without establishing population generalisation or width accuracy. The method therefore offers an inspectable, deterministic route to projected filament identity; broader claims require independent real fields, validated physical widths, further released implementations run in a matched comparison, and an external-generator benchmark with sealed seeds.

CRediT author contribution statement

[To do: Fill in each author’s CRediT roles before submission, checked against actual contributions; keep the author names and order as listed above.]

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data and code availability

The PLECTA source, complete fixed configuration, compact seed manifest, and aggregate machine-readable evidence are available at <https://github.com/Jorallan/PLECTA>. Generated scenes, detailed per-scene records, and development diagnostics are retained locally. **[To do: Create a tagged archival release and add its DOI before submission; do not infer or preassign a DOI.]**

Acknowledgements

The authors thank the maintainers of NumPy, SciPy, scikit-image, NetworkX, and Matplotlib. Pipeline A, the upstream synthetic-data-driven mask workflow that supplies the masks for the SEM demonstration, adapts the open CycleGAN/U-Net particle-segmentation workflow of Rühle et al. [32] and carries its own evaluation elsewhere.

References

- [1] Michael F. L. De Volder, Sameh H. Tawfick, Ray H. Baughman, and A. John Hart. Carbon nanotubes: Present and future commercial applications. *Science*, 339(6119):535–539, 2013. doi: 10.1126/science.1222453.
- [2] Jonathan N. Coleman, Umar Khan, Werner J. Blau, and Yurii K. Gun’ko. Small but strong: A review of the mechanical properties of carbon nanotube-polymer composites. *Carbon*, 44(9):1624–1652, 2006. doi: 10.1016/j.carbon.2006.02.038.
- [3] Alexey N. Volkov and Leonid V. Zhigilei. Structural stability of carbon nanotube films: The role of bending buckling. *ACS Nano*, 4(10):6187–6195, 2010. doi: 10.1021/nn1015902.
- [4] Igor Ostanin, Roberto Ballarini, David Potyondy, and Traian Dumitrică. A distinct element method for large scale simulations of carbon nanotube assemblies. *Journal of the Mechanics and Physics of Solids*, 61(3):762–782, 2013. doi: 10.1016/j.jmps.2012.10.016.
- [5] Bernard K. Wittmaack, Abu Horaira Banna, Alexey N. Volkov, and Leonid V. Zhigilei. Mesoscopic modeling of structural self-organization of carbon nanotubes into vertically aligned networks of nanotube bundles. *Carbon*, 130:69–86, 2018. doi: 10.1016/j.carbon.2017.12.078.
- [6] Liyu Dong, Haibin Hang, Jin Gyu Park, Washington Mio, and Richard Liang. Detecting carbon nanotube orientation with topological analysis of scanning electron micrographs. *Nanomaterials*, 12(8):1251, 2022. doi: 10.3390/nano12081251.
- [7] Tamjid Imtiaz, Jacques Doumani, Fuyang Tay, Natsumi Komatsu, Stephen Marcon, Motonori Nakamura, Saunab Ghosh, Andrey Baydin, Junichiro Kono, and Ahmed Zubair. Facile alignment estimation in carbon nanotube films using image processing. *Signal Processing*, 202:108751, 2023. doi: 10.1016/j.sigpro.2022.108751.
- [8] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-Net: Convolutional networks for biomedical image segmentation. In *Medical Image Computing and Computer-Assisted Intervention – MICCAI 2015*, volume 9351 of *Lecture Notes in Computer Science*, pages 234–241, 2015. doi: 10.1007/978-3-319-24574-4_28.
- [9] Fabian Isensee, Paul F. Jaeger, Simon A. A. Kohl, Jens Petersen, and Klaus H. Maier-Hein. nnU-Net: A self-configuring method for deep learning-based biomedical image segmentation. *Nature Methods*, 18:203–211, 2021. doi: 10.1038/s41592-020-01008-z.
- [10] Carsen Stringer, Tim Wang, Michalis Michaelos, and Marius Pachitariu. Cellpose: A generalist algorithm for cellular segmentation. *Nature Methods*, 18:100–106, 2021. doi: 10.1038/s41592-020-01018-x.
- [11] Uwe Schmidt, Martin Weigert, Coleman Broaddus, and Gene Myers. Cell detection with star-convex polygons. In *Medical Image Computing and Computer Assisted Intervention – MICCAI 2018*, *Lecture Notes in Computer Science*, pages 265–273, 2018. doi: 10.1007/978-3-030-00934-2_30.

- [12] Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, Tete Xiao, Spencer Whitehead, Alexander C. Berg, Wan-Yen Lo, Piotr Dollár, and Ross Girshick. Segment anything. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 3992–4003, 2023. doi: 10.1109/ICCV51070.2023.00371.
- [13] Carsten Steger. An unbiased detector of curvilinear structures. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 20(2):113–125, 1998. doi: 10.1109/34.659930.
- [14] Alejandro F. Frangi, Wiro J. Niessen, Koen L. Vincken, and Max A. Viergever. Multiscale vessel enhancement filtering. In *Medical Image Computing and Computer-Assisted Intervention – MICCAI’98*, volume 1496 of *Lecture Notes in Computer Science*, pages 130–137, 1998. doi: 10.1007/BFb0056195.
- [15] Yoshinobu Sato, Shin Nakajima, Nobuyuki Shiraga, Hideki Atsumi, Shigeyuki Yoshida, Thomas Koller, Guido Gerig, and Ron Kikinis. Three-dimensional multi-scale line filter for segmentation and visualization of curvilinear structures in medical images. *Medical Image Analysis*, 2(2):143–168, 1998. doi: 10.1016/S1361-8415(98)80009-1.
- [16] Andrew M. Stein, David A. Vader, Louise M. Jawerth, David A. Weitz, and Leonard M. Sander. An algorithm for extracting the network geometry of three-dimensional collagen gels. *Journal of Microscopy*, 232(3):463–475, 2008. doi: 10.1111/j.1365-2818.2008.02141.x.
- [17] Jeremy S. Bredfeldt, Yuming Liu, Carolyn A. Pehlke, Matthew W. Conklin, Joseph M. Szulcowski, David R. Inman, Patricia J. Keely, Robert D. Nowak, Thomas R. Mackie, and Kevin W. Eliceiri. Computational segmentation of collagen fibers from second-harmonic generation images of breast cancer. *Journal of Biomedical Optics*, 19(1):016007, 2014. doi: 10.1117/1.JBO.19.1.016007.
- [18] Ting Xu, Dimitrios Vavylonis, Feng-Ching Tsai, Gijsje H. Koenderink, Wei Nie, Eddy Yusuf, I-Ju Lee, Jian-Qiu Wu, and Xiaolei Huang. SOAX: A software for quantification of 3D biopolymer networks. *Scientific Reports*, 5:9081, 2015. doi: 10.1038/srep09081.
- [19] Daniel A. D. Flormann, Moritz Schu, Emmanuel Terriac, Divyendu Thalla, Lucina Kainka, Marcus Koch, Annica K. B. Gad, and Franziska Lautenschläger. A novel universal algorithm for filament network tracing and cytoskeleton analysis. *The FASEB Journal*, 35(5):e21582, 2021. doi: 10.1096/fj.202100048R.
- [20] Nathan A. Hotaling, Kapil Bharti, Haydn Kriel, and Carl G. Simon. DiameterJ: A validated open source nanofiber diameter measurement tool. *Biomaterials*, 61:327–338, 2015. doi: 10.1016/j.biomaterials.2015.05.015.
- [21] Saurav Basu, Chi Liu, and Gustavo Kunde Rohde. Extraction of individual filaments from 2D confocal microscopy images of flat cells. *IEEE/ACM Transactions on Computational Biology and Bioinformatics*, 12(3):632–643, 2015. doi: 10.1109/TCBB.2014.2372783.
- [22] David Breuer and Zoran Nikoloski. DeFiNe: An optimisation-based method for robust disentangling of filamentous networks. *Scientific Reports*, 5:18267, 2015. doi: 10.1038/srep18267.
- [23] Zhen Zhang, Yukako Nishimura, and Pakorn Kanchanawong. Extracting microtubule networks from superresolution single-molecule localization microscopy data. *Molecular Biology of the Cell*, 28(2):333–345, 2017. doi: 10.1091/mbc.E16-06-0421.
- [24] Jaydeep De, Li Cheng, Xiaowei Zhang, Feng Lin, Huiqi Li, Kok Haur Ong, Weimiao Yu, Yuanhong Yu, and Sohail Ahmed. A graph-theoretical approach for tracing filamentary structures in neuronal and retinal images. *IEEE Transactions on Medical Imaging*, 35(1): 257–272, 2016. doi: 10.1109/TMI.2015.2465962.

- [25] Viktor Wegmayr, Aytunc Sahin, Björn Sæmundsson, and Joachim M. Buhmann. Instance segmentation for the quantification of microplastic fiber images. In *2020 IEEE Winter Conference on Applications of Computer Vision (WACV)*, pages 2199–2206, 2020. doi: 10.1109/WACV45572.2020.9093352.
- [26] Clément Ployat, Yosra Mehrjoo, Renaud Duval, Marie Carole Boucher, Santiago Costantino, and Hugo Wurtele. DNAi: An open-source AI tool for unbiased DNA fiber analysis. *Nucleic Acids Research*, 54(7):gkag335, 2026. doi: 10.1093/nar/gkag335.
- [27] Isabella Østerlund, Arne Neumann, Zhiming Ma, Yansong Miao, Staffan Persson, and Zoran Nikoloski. GraFT: A robust network-based spatiotemporal analysis of filamentous structures. *Science Advances*, 12(13):eadz4132, 2026. doi: 10.1126/sciadv.adz4132.
- [28] Yi Liu, Abhishek Kolagunda, Wayne Treible, Alex Nedo, Jeffrey Caplan, and Chandra Kambhamettu. Intersection to overpass: Instance segmentation on filamentous structures with an orientation-aware neural network and terminus pairing algorithm. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 125–133, 2019. doi: 10.1109/CVPRW.2019.00021.
- [29] Max Frei and Frank Einar Kruis. FibeR-CNN: Expanding Mask R-CNN to improve image-based fiber analysis. *Powder Technology*, 377:974–991, 2021. doi: 10.1016/j.powtec.2020.08.034. Published online 2020.
- [30] Yi Liu, Su Peng, Jeffrey Caplan, and Chandra Kambhamettu. Pick and trace: Instance segmentation for filamentous objects with a recurrent neural network. In *Medical Image Computing and Computer Assisted Intervention – MICCAI 2023*, Lecture Notes in Computer Science, pages 635–645, 2023. doi: 10.1007/978-3-031-43993-3_61.
- [31] Torben Peters, John Schumann, Kerstin Kämpf, and Asmus Meyer-Plath. AI-assisted recognition of fibre instances in scanning electron microscope images for aerosol characterization and exposure assessment. *Powder Technology*, 469:121909, 2026. doi: 10.1016/j.powtec.2025.121909.
- [32] Bastian Rühle, Julian Frederic Krumrey, and Vasile-Dan Hodoroaba. Workflow towards automated segmentation of agglomerated, non-spherical particles from electron microscopy images using artificial neural networks. *Scientific Reports*, 11(1):4942, 2021. doi: 10.1038/s41598-021-84287-6.
- [33] Oday Allan. CNT_SEM: CycleGAN synthesis and U-Net segmentation workflows for carbon-nanotube SEM data. GitHub repository, 2026. URL https://github.com/Jorallan/CNT_SEM.
- [34] William M. Rand. Objective criteria for the evaluation of clustering methods. *Journal of the American Statistical Association*, 66(336):846–850, 1971. doi: 10.1080/01621459.1971.10482356.
- [35] Ignacio Arganda-Carreras, Srinivas C. Turaga, Daniel R. Berger, Dan Cireşan, Alessandro Giusti, Luca M. Gambardella, Jürgen Schmidhuber, Dmitry Laptev, Sarvesh Dwivedi, Joachim M. Buhmann, Ting Liu, Mojtaba Seyedhosseini, Tolga Tasdizen, Lee Kamentsky, Radim Burget, Václav Uher, Xiao Tan, Changming Sun, Tuan D. Pham, Erhan Bas, Mustafa G. Uzunbas, Albert Cardona, Johannes Schindelin, and H. Sebastian Seung. Crowdsourcing the creation of image segmentation algorithms for connectomics. *Frontiers in Neuroanatomy*, 9:142, 2015. doi: 10.3389/fnana.2015.00142.
- [36] Lawrence Hubert and Phipps Arabie. Comparing partitions. *Journal of Classification*, 2(1): 193–218, 1985. doi: 10.1007/BF01908075.

- [37] Marina Meilă. Comparing clusterings—an information based distance. *Journal of Multivariate Analysis*, 98(5):873–895, 2007. doi: 10.1016/j.jmva.2006.11.013.