

PLECTA: Overlap-aware instance reconstruction of dense filament networks from binary masks

Oday Allan^a Stefan Luding^a Igor Ostanin^{a,*}

^a Multi-Scale Mechanics (MSM), Faculty of Engineering Technology, University of Twente, Enschede, The Netherlands

[To do: Add ORCIDs.]

Abstract

Dense filament networks are difficult to partition into instances because crossings join distinct objects while segmentation gaps divide single objects. We present PLECTA, a deterministic method that converts a binary filament-axis mask into overlap-aware two-dimensional instances. PLECTA skeletonizes and decomposes the mask into arms and junctions, alternates exact per-junction continuation matching with globally gated gap matching, and re-estimates tangent, chord, and curvature evidence from the assembled chains. Crossing pixels may therefore belong to more than one output instance.

Evaluation used a method-independent partition of each input skeleton into common fragments. On 128 independently seeded held-out scenes from the same procedural generator, PLECTA achieved mean pairwise $F_1=0.854$ (density-stratified 95% bootstrap CI [0.845, 0.864]), precision 0.868, recall 0.843, recovery 0.799, and adjusted Rand index 0.851, with no failed scenes. Mean F_1 decreased from 0.911 at 20% coverage to 0.775 at 60%, identifying dense crossings as the main remaining regime of error. Development-only ablations showed that chord-turn evidence, gap bridging, and junction-topology consolidation made the largest contributions.

CNT-sheet scanning electron microscopy provides the application case. An optional upstream synthetic-data-driven U-Net supplies the binary mask; SEM width measurement and ribbon rendering are downstream characterisation stages and did not affect the held-out grouping result. The evidence supports same-generator synthetic generalisation and a development illustration on real SEM data, not general real-image or three-dimensional reconstruction claims.

Keywords: filament networks; instance segmentation; graph matching; geometric reconstruction; carbon nanotubes; scanning electron microscopy.

1 Introduction

The properties of carbon-nanotube (CNT) sheets, yarns, and composites depend on the geometry of their bundle networks: length, width, orientation, curvature, and contact topology influence mechanical and electrical response [1–5]. Scanning electron microscopy (SEM) resolves this projected network, but routine image analysis often stops at coverage or orientation distributions [6, 7]. Instance-resolved analysis instead requires each visible bundle to remain addressable across gaps, touches, and crossings.

This requirement is distinct from semantic segmentation. U-Net and its self-configuring variants can identify filament foreground [8, 9], but a binary foreground mask does not encode object identity. A faded segment splits one filament into several connected components; a crossing joins different filaments into one component. The ambiguity is greatest in dense networks, where several local continuations can be geometrically plausible. Compact-object instance

*Corresponding author: i.ostanin@utwente.nl

priors—bounded blobs, centre-directed flows, or star-convex shapes—are poorly matched to long open curves that may overlap [10–12].

Existing methods address complementary parts of the problem. Ridge and vesselness filters enhance curvilinear signal but do not assign identity across junctions [13–15]. FIRE, CT-FIRE, SOAX, FiNTA, and DiameterJ extract fibres or network statistics in collagen, biopolymer, and fibrous materials [16–20]; methods such as that of Depriester et al. address fibre separation directly [21].

The reconstruction of individual filaments from networks is itself established prior art. Basu et al. combine centreline fragments by integer programming, DeFiNe optimizes filament covers of weighted graphs, and De et al. partition directed graphs to resolve crossovers [22–24]. SIFNE removes junctions and reconnects termini using propagation direction, angular continuity, and gap distance [25]; DNAi evaluates weighted local endpoint pairings after erasing fibre junctions [26]. Learned alternatives include orientation-aware terminus pairing, deep pixel embeddings for tangled microplastic fibres, FibeR-CNN, and recurrent pick-and-trace models [27–30]. GraFT additionally traces and tracks actin structures as a graph over time [31]. Thus skeleton graphs, continuation angles, distance gates, optimization, and filament-level outputs are not claimed here as new in isolation.

CNT-specific workflows increasingly combine semantic segmentation with algorithmic tracing and transverse intensity profiles [32]. The present work focuses on the difficult middle step of that workflow: binary mask to overlap-aware instances. This separation is deliberate. It allows grouping to be evaluated independently of the upstream segmenter and prevents width rendering from changing an identity metric. An optional CycleGAN/U-Net route supplies masks for the SEM application [33, 34], but it is not the principal contribution.

We introduce PLECTA (Pairwise Linking with Evidence from Chain Tangents and Assignment), a deterministic graph-matching method for dense filament networks. Its contribution is the specific binary-mask-only reconstruction design evaluated here: graph preconditioning; exact partial matching within each merged junction, including the option to leave an arm unmatched; global gated matching across disconnected gaps; iterative replacement of local tangents by chain-supported frames; and overlap-aware output layers in which crossing pixels may be shared. The name evokes the Latin *plectere*, to weave or interlace, without implying recovery of depth.


The study asks three questions. First, how accurately does the fixed method group method-independent skeleton fragments on unseen procedural scenes, including high-density cases? Second, which active geometric and topological components account for that performance, and how sensitive is it to mask degradation? Third, what does the method deliver in a CNT SEM workflow when the mask is derived manually or by a learned segmenter? Synthetic scenes provide exact overlap-aware identity; real SEM remains a development case. The intended claim is therefore a reproducible instance-reconstruction result under the tested generator, with explicit limits on distribution transfer, width validation, and 3-D interpretation.

2 Materials and methods

2.1 Scope and workflow

The formal input is a binary filament-axis mask $M \in \{0, 1\}^{H \times W}$. PLECTA reconstructs an overlap-aware collection $\{I_1, \dots, I_N\}$, where $I_n \subseteq M$ and a crossing pixel may belong to several instances. The evaluated grouping stage uses only M . A grayscale SEM image is optional downstream input for width and brightness measurement; it cannot revise the fixed grouping. Figure 1 separates this core task from upstream mask extraction and downstream rendering.

Input-mask formation



Only the green core contributed to the 128-scene held-out grouping result.

Figure 1. Workflow and evidence boundary. A micrograph is converted to a binary axis mask manually or by an upstream segmenter. PLECTA groups the mask into overlap-aware 2-D instances. SEM characterisation and ribbon rendering are optional downstream operations and were absent from the held-out grouping result.

2.2 Synthetic data and evaluation split

Procedural scenes preserve the ordered centreline $\mathcal{C}_k$ and rendered support Ω_k of every filament k . Individual supports remain separate, so several filaments may own the same crossing pixel. The generator provides an undegraded one-pixel axis mask, a degraded thin mask containing small gaps and irregularities, a combined thick foreground, an SEM-like image, and the exact overlap-aware instance reference (Figure 2). Areal coverage is the union area of the thick supports divided by image area; it is not the occupancy of the thin input mask.

All evaluated scenes were 512×512 px. The recorded generator uses 2-px centreline steps and Gaussian turn increments with standard deviation 0.035 rad, smoothed over 20 px. Path length is sampled uniformly from 0.25–0.85 of the image diagonal; the anchor and initial direction are uniform, and the longest contiguous in-frame part is retained. Apparent bundle width and brightness are sampled uniformly from 7–16 px and 0.55–0.98. Filaments are added until the union of their thick supports reaches the requested coverage. The degraded axis is obtained by thresholding a longitudinally modulated thin ridge at 0.27 after correlated edge noise (amplitude 0.06), then applying recorded hard-cut probability 0.12, crossing-occlusion probability 0.15, a Poisson mean of 2.5 spurious marks, and light morphological roughening. The primary reconstruction input is the one-pixel version; clean and 2- and 3-px variants are retained as controlled derivatives of the same geometry and degradation draw. These are generator settings, not fitted physical distributions of CNT morphology.

Development used 84 scenes: 20 initial scenes across 20–60% coverage, 48 additional scenes at 20, 30, and 60%, and 16 scenes at 40%. The held-out set comprised 128 different seeds: 32 scenes at each of 20, 30, 40, and 60% coverage. It is an independent sample from the same generator, not an independent distribution. A separate 125-scene locked set remains unevaluated. The compact seed manifest records every seed block and an executable verifier checks counts, disjointness, and available scene metadata.

A separate controlled robustness set contained 50 scenes, ten at each of 20, 30, 40, 50, and 60% coverage. Each geometry was scored once from its clean axis and once from its degraded mask. This set was used only to quantify input-degradation sensitivity after the configuration was fixed.

Thick-bundle / thin-axis synthetic scenes with exact ground truth




Figure 2. Procedural representations derived from the same filament geometry. Instance identity is retained through crossings; the degraded thin-axis mask is the primary algorithm input.

2.3 CNT SEM data and mask sources

Two manually annotated CNT-sheet SEM fields were available for exploratory analysis: B58_100 (1024×1536 px, 280 filaments) and B58_110 (1024×1536 px, 430 filaments). The nominal horizontal field width is $1.66 \mu\text{m}$; project metadata associate it with 1536 px, or approximately 1.08 nm px^{-1} . The annotation is one human interpretation of projected, locally ambiguous crossings. Both fields informed development and neither is a population-level test.

Two inputs were derived per field. The first was the skeleton of the manual instance annotation and is therefore an oracle-input condition, not an independent segmentation. The second was predicted by an upstream nnU-Net trained on paired synthetic images and masks; CycleGAN translation is one route from procedural masks to SEM-like training images [33, 34]. B58_100 appeared in that model’s training data, whereas B58_110 was verified absent from the relevant training set. Consequently, only B58_110 provides a training-clean observation of mask transfer, and $n = 2$ precludes inference.

Mask quality was compared after skeletonisation with a symmetric 3-px tolerance, avoiding a reward for mask thickness. Reconstruction scores were computed separately within each mask’s visible common-fragment population; because those populations differ, cross-mask F_1 values are descriptive rather than paired estimates.

Author metadata required before submission: SEM instrument and detector, accelerating voltage, working distance, specimen preparation, CNT material identity, annotation protocol, and confirmation of the 1536-pixel calibration against the stored image dimensions. The upstream model archive must also supply its architecture, training split, and optimisation settings.

2.4 PLECTA reconstruction

2.4.1 Skeleton graph and local frames

PLECTA binarises M at > 0 , skeletonises it, and iteratively removes spurs no longer than 3 px. Skeleton pixels with degree at least three form junction clusters, and adjacent clusters are




Figure 3. Geometric evidence and exact pairing. (a) An arclength fit over a fixed support window gives the outward tangent and signed curvature at each stub. (b) Candidate cost combines direct tangent reversal θ , the two turns ϕ_a, ϕ_b onto the tip-to-tip chord, distance for gap candidates, and curvature continuity; the chord contribution fades at short separation. (c) All feasible pairs at one junction are scored jointly. Exact partial matching selects disjoint continuations of minimum total cost while allowing a stub to remain unmatched at a fixed price.

merged into nodes. Connectors up to 5 px are absorbed into one node; connectors up to 14 px merge nearby nodes but remain real arms; free debris up to 2 px is absorbed. Each remaining connected path is stored as an ordered arm with one stub at each end.

For ordered points $\mathbf{x}(s) = (r(s), c(s))$ running from a stub tip into its arm, PLECTA fits each coordinate by least squares in arclength $0 \leq s \leq w$. The fit is quadratic when the observed span is at least 18 px and linear otherwise. With $\det(\mathbf{a}, \mathbf{b}) = a_r b_c - a_c b_r$, the outward frame is

$$\mathbf{p} = \mathbf{x}(0), \quad \mathbf{t} = -\frac{\mathbf{x}'(0)}{\|\mathbf{x}'(0)\|}, \quad \kappa = \frac{\det[-\mathbf{x}'(0), \mathbf{x}''(0)]}{\|\mathbf{x}'(0)\|^3}. \quad (1)$$

A frame is reliable only when at least three samples span at least 3 px. The initial fit uses $w = 24$ px. Subsequent fits may walk through the preceding round's links and use $w = 55$ px, allowing several short arms to provide one longer directional estimate.

For stubs a, b , let $d = \|\mathbf{p}_b - \mathbf{p}_a\|$ and, when $d > 0$, $\mathbf{u} = (\mathbf{p}_b - \mathbf{p}_a)/d$. The direct reversal, two chord turns, and chord-confidence factor are

$$\theta = \cos^{-1}(\mathbf{t}_a \cdot -\mathbf{t}_b), \quad \phi_a = \cos^{-1}(\mathbf{t}_a \cdot \mathbf{u}), \quad \phi_b = \cos^{-1}(\mathbf{t}_b \cdot -\mathbf{u}), \quad q(d) = \min(1, d/d_0). \quad (2)$$

$$C_{ab} = w_d \theta + w_t q(d)(\phi_a + \phi_b) + w_\ell d/\ell_0 + 30w_\kappa |\kappa_a + \kappa_b|. \quad (3)$$

For coincident tips the two chord angles are set to zero. Junction costs set $w_\ell = 0$; gap costs retain the length term. Thus the tangent-reversal term remains usable at a rasterised crossing, while $q(d)$ suppresses the poorly determined chord direction until its length reaches d_0 .

2.4.2 Candidate gates and exact matching

Let S be the stubs considered at one junction, or all free stubs in the gap stage; let E contain the geometrically admissible pairs; and let p_i be the price for leaving stub i unmatched. PLECTA solves the exact partial matching problem

$$M^* = \arg \min_{M \subseteq E: M \text{ is a matching}} \left[\sum_{(i,j) \in M} C_{ij} + \sum_{i \in S \setminus V(M)} p_i \right]. \quad (4)$$

Table 1. Complete materialised PLECTA parameter configuration used for the reported reconstruction. Angles are in radians and lengths in pixels; fixed implementation rules are stated in the surrounding text.

Group	Parameters	Values
Graph	spur; absorbed bridge; free debris; node join	3; 5; 2; 14
Frame	local window; chain window; quadratic span	24; 55; 18
Junction	w_d ; w_t ; chord floor d_0 ; w_κ ; unmatched price; gap relief	0.6; 1.2; 4; 1.0; 0.62; 0
Gap	maximum length; w_d ; w_t ; w_ℓ ; length scale ℓ_0 ; chord floor d_0 ; w_κ ; unmatched price; maximum θ ; maximum ϕ_a, ϕ_b	85; 0.2; 1.0; 0.15; 60; 3; 0.5; 0.3; 0.40; 0.28
Rounds	number; initial strictness	8; 0.85

It is implemented as unconstrained-cardinality maximum-weight matching with edge weight $p_i + p_j - C_{ij}$; edges with $C_{ij} \geq p_i + p_j$ are omitted. Every junction is solved independently, with candidates restricted to different arms sharing that node. At full strictness $p_i = 0.62$, so a junction edge requires $C_{ij} < 1.24$. The configured gap-relief term is zero, hence all junction stubs have the same unmatched price.

Gap candidates are constructed only from still-unmatched, reliable stubs. They must belong to different arms and different junctions, lie within 85 px, have direct tangent mismatch at most 0.40 rad, and have chord-turn mismatch at each end at most 0.28 rad. Global exact matching then accepts only edges whose cost is below 0.60 at full strictness.

For rounds $r = 0, \dots, 7$, the strictness factor is $\alpha_r = 0.85 + 0.15r/7$. Only the junction and gap unmatched prices and the two gap-angle limits are multiplied by α_r ; distance and all cost weights remain fixed. Each round estimates frames from the preceding matching, remakes all junction links, removes cycles, remakes all gap links, and removes cycles again. In the arm-link graph, a component with n arms closes a cycle when it contains at least n accepted links; PLECTA repeatedly deletes that component’s matched pair with the greatest tip-to-tip distance until it is open. No cycle was encountered in the 84 development scenes. Only $r = 7$ has $\alpha_r = 1$, so its full-strictness matching is returned; the objective does not select an earlier, deliberately under-linked annealing round.


2.4.3 Instance output

Connected components of accepted arm links define instances. PLECTA paints every arm pixel and every junction cluster touched by a chain into that instance, so intersecting instances share crossing pixels. An isolated one-arm component shorter than 6 px is omitted. Accepted gap links determine identity but are not painted into the fixed centreline output. All effective parameters, including values previously inherited as code defaults, are materialised in the fixed JSON configuration.

2.5 Optional SEM characterisation and ribbon rendering

After grouping, perpendicular SEM profiles are sampled every 2 px along each ordered chain and at least 6 px from a junction. Profiles are recentred by at most 3 px and accepted only when contrast exceeds both four background-noise standard deviations and 0.05 normalised intensity; fitted widths must lie between 1.5 and 36 px. At least three valid cuts are required. These measurements attach apparent width, brightness, uncertainty, and a quality flag to an existing instance; they cannot alter its arm membership.

Optional ribbon rendering interpolates accepted gaps, resamples the centreline at 1 px, smooths it over 9 px with fixed endpoints, and draws the chain at its fitted width clipped to 3–40 px.



Gap links affect identity but are not painted into the core centreline output.

Figure 4. PLECTA stage and gate sequence. Topology preconditioning produces arms, stubs, and junction nodes. Every round applies exact per-junction matching, a cycle guard, global gated gap matching, a second cycle guard, and chain-supported frame re-estimation. The final chains are painted as overlap-aware instance layers.

It excludes junction blobs and does not clip the ribbon back to the input mask. Silhouette absorption would union rendered instances, but its threshold is zero by default. No absorption or later post-hoc continuation merge contributed to the 128-scene held-out result; neither is presented as part of the evaluated algorithm.

Synthetic width recovery was assessed only on the 84-scene development set. A measured chain was associated with a rendered reference filament when at least 90% of its sampled centreline belonged to that filament. Apparent width was then compared with the generator value for the associated filament; interval calibration used scene-level cross-validation so that cuts from one scene could not occur in both calibration and assessment folds. This validates measurement of the synthetic renderer, not physical CNT diameter.

2.6 Metrics and statistical analysis

The primary endpoint is common-fragment pairwise F_1 . For each input mask, the scorer skeletonises once, iteratively prunes junction-touching free-end spurs of at most 3 px, removes pixels with at least three 8-neighbours, labels the remaining 8-connected pieces, and discards pieces shorter than 6 px. This partition is independent of PLECTA and of the reference instances.

Each retained fragment is assigned separately to the overlap-aware reference layers and prediction layers. The instance with the largest direct pixel overlap is selected if it covers at least 30% of the fragment. Otherwise, instance masks are flattened deterministically in decreasing area and then increasing identifier order; the fragment is assigned to the most frequent nearest flattened label among its pixels lying within 6 px of an instance. Ties resolve to the smaller identifier. If neither rule finds a candidate, the fragment is unassigned.

Let n_{ij} be the number of fragments assigned to reference group i and predicted group j , with $a_i = \sum_j n_{ij}$ and $b_j = \sum_i n_{ij}$. Every unassigned fragment is treated as its own singleton rather than placing all missing evidence in one artificial group. The pair counts are

$$\text{TP} = \sum_{ij} \binom{n_{ij}}{2}, \quad \text{FP} = \sum_j \binom{b_j}{2} - \text{TP}, \quad \text{FN} = \sum_i \binom{a_i}{2} - \text{TP}. \quad (5)$$

A true-positive pair therefore shares both identities; a false-positive pair joins different reference filaments; and a false-negative pair separates fragments of one reference filament. Thus

$$P = \frac{\text{TP}}{\text{TP} + \text{FP}}, \quad R = \frac{\text{TP}}{\text{TP} + \text{FN}}, \quad F_1 = \frac{2PR}{P + R}. \quad (6)$$

This is a clustering metric related to the Rand family [35, 36]. It scores grouping, not rendered width, and weights filaments with many fragments more strongly because their pair count is quadratic. We therefore also report adjusted Rand index [37], variation of information split and merge components [38], and the fraction of reference filaments recovered by one dominant predicted instance using 0.8 purity and coverage thresholds.

Scores are computed per scene and reported summaries are arithmetic means over successfully scored scenes, rather than one pooled contingency table. A scene with fewer than two fragments, or with no positive reference or predicted pairs, has no pair evidence and its undefined precision or recall is retained as missing rather than replaced by a favourable value. The evaluation harness writes an explicit error row for any failed scene; such rows are counted as failures and excluded from metric means. The reported 128-scene held-out run had no failures.

A diagnostic connectivity floor assigns one label to each 8-connected component of the same held-out input mask and is scored without tuning by the same common-fragment procedure. It tests whether the benchmark can be solved by connectivity alone; it is not treated as a competitive filament tracer.

Held-out means and density strata are accompanied by 95% nonparametric bootstrap intervals (20,000 replicates; seed 20,260,810). Aggregate resampling is stratified by density. The clean-minus-degraded analysis resamples paired scene differences within density. Active-component ablations are one-at-a-time development experiments on all 84 development scenes and are labelled as such; they do not reuse the held-out set.

3 Results

3.1 Held-out synthetic reconstruction

PLECTA completed all 128 held-out scenes with 0 failures. Mean common-fragment pairwise F_1 was 0.854 (density-stratified 95% bootstrap CI [0.845, 0.864]); mean precision, recall, reference-instance recovery, and adjusted Rand index were 0.868, 0.843, 0.799, and 0.851, respectively (Table 2). These scenes used unseen seeds but came from the same procedural generator as development data, so the result measures same-generator generalisation rather than distribution transfer.

As a diagnostic floor, assigning one instance label to every connected component of the same held-out masks produced $F_1 = 0.085$ (precision 0.047, recall 0.850, ARI 0.035). Its high recall but very low precision and VI-merge of 4.103 show the expected failure: crossings collapse many distinct filaments into a single component. This is not a competitive tracer, but it verifies that connectivity alone cannot solve the evaluated grouping task.

Performance decreased with scene density (Figure 6a). Mean F_1 fell from 0.911 [0.888, 0.933] at 20% coverage to 0.775 [0.760, 0.790] at 60% coverage; the intermediate 30% and 40% means were

Table 2. Held-out reconstruction scores overall and by target areal coverage. VI denotes variation of information; lower VI is better, whereas higher values are better for the other metrics.

Group	n	F_1	Precision	Recall	Recovery	ARI	VI split	VI merge
All	128	0.854	0.868	0.843	0.799	0.851	0.256	0.216
20%	32	0.911	0.928	0.896	0.918	0.907	0.137	0.093
30%	32	0.874	0.892	0.858	0.828	0.870	0.211	0.161
40%	32	0.858	0.864	0.853	0.796	0.855	0.242	0.224
60%	32	0.775	0.787	0.765	0.656	0.772	0.436	0.388




Figure 5. Rule-selected held-out examples at 20% (top) and 60% (bottom) target coverage. For each stratum, the lexicographically first scene identifier was used; neither score nor appearance entered selection. From left to right: degraded input mask, overlap-aware reference layers, PLECTA layers, and common-fragment grouping. Green fragments belong to a one-to-one matched group; red marks a split, merge, or missing prediction; purple marks a fragment unassigned in the reference. Dark pixels with a yellow halo belong to more than one instance. The manifest records scene identifiers, seeds, scores, and the parameter-file hash.

0.874 and 0.858. Recorded per-scene runtime had mean 2.54 s, median 1.12 s, 90th percentile 6.42 s, and maximum 34.06 s; these timings are hardware-specific.

Figure 5 shows the corresponding grouping behaviour rather than only aggregate scores. Most fragments are assigned consistently in the sparse scene. In the dense scene, the longer red runs and reduced recovery expose both erroneous unions and interrupted identities around crowded junctions; the prediction nevertheless retains explicit overlap at crossings instead of forcing a single flat label.

3.2 Input-mask robustness

On the separate 50-scene controlled set, mean F_1 was 0.822 from degraded axes and 0.844 from the corresponding clean axes. The paired clean-minus-degraded change was 0.0224 (95% CI [0.0113, 0.0333]). Thus the modelled gaps and raster irregularities caused a small resolved loss under this generator, without testing robustness to a different image distribution.

Table 3. One-at-a-time component ablations on the 84-scene development set. Negative ΔF_1 values denote losses relative to the evaluated configuration.

Development variant	F_1	ΔF_1	ARI
Fixed configuration	0.867	+0.000	0.864
No Gap Bridging	0.728	-0.139	0.722
Single Round	0.847	-0.020	0.844
No Chain-Extended Frames	0.846	-0.021	0.842
No Annealing	0.863	-0.004	0.860
No Curvature Term	0.863	-0.005	0.859
No Join_Px	0.845	-0.022	0.841
No Junction-Cluster Consolidation	0.804	-0.063	0.799
No Junction Chord-Turn	0.681	-0.186	0.675




Figure 6. PLECTA performance across target areal coverage and component ablations. (a) Held-out mean common-fragment F_1 with 95% bootstrap intervals (32 scenes per level). (b) Development-only changes from selected one-at-a-time ablations; more negative values indicate larger losses.

3.3 Development ablations and downstream rendering

One-at-a-time ablations were evaluated only on the 84 development scenes (Table 3; Figure 6b). Relative to the evaluated configuration, removing the junction chord-turn term changed F_1 by -0.1858, removing gap matching by -0.1394, and disabling connector- based junction-cluster consolidation by -0.0634. Removing chain-extended frames and using a single round produced smaller changes of -0.0215 and -0.0199. These development results identify chord-aware continuation, gap matching, and junction-topology preconditioning as the largest tested contributors; they are not held-out effect estimates.

Default Stage 4 ribbon rendering achieved development $F_1 = 0.860$, a change of -0.0068 from the core centreline output. The rendered instances had disconnected-component fraction 0.000 and branched fraction 0.005. Rendering therefore improved mask continuity while slightly changing the representation score, without changing chain membership. Enabling the separate 0.6 silhouette-absorption rule reduced F_1 further to 0.857. Because absorption can merge identities, it remained disabled and neither rendering condition contributed to the held-out grouping result. These topology measurements do not validate apparent width.

The optional image-characterisation stage was separately checked against the known synthetic renderer on 84 development scenes. Among 3059 instances matched to one reference filament

Table 4. Exploratory reconstruction from manual-derived and U-Net axes in two CNT SEM fields. Axis recall and precision compare the U-Net skeleton with the manual-derived skeleton using a symmetric 3-px tolerance.

Field	Input mask	Axis recall	Axis precision	F_1	ARI
B58_100	Manual-derived	—	—	0.882	0.882
B58_100	U-Net	0.782	0.753	0.438	0.435
B58_110	Manual-derived	—	—	0.837	0.836
B58_110	U-Net	0.729	0.887	0.452	0.449




Figure 7. Training-clean B58_110 mask-source case. Left: the manual-instance-derived axis. Centre: the upstream U-Net axis. Right: tolerance-matched manual axis (green), missed manual axis (red), and U-Net axis without a matched manual axis (blue). The comparison illustrates upstream omissions and fragmentation; it is not a PLECTA output panel.

with purity at least 0.9, fitted apparent width had Pearson $r = 0.9638$, mean absolute error 0.268 px, and median absolute error 0.132 px. A quality-stratified 90% interval, calibrated by leaving scenes out, covered 0.899 of instances. This validates recovery of the synthetic renderer’s width under the development model; it is not a validation of physical CNT diameter.

3.4 Two-field CNT SEM mask-source case

The two real SEM fields were exploratory development cases (Table 4). From the manual-derived axis, PLECTA reached $F_1=0.882$ on B58_100 and 0.837 on B58_110. These are oracle-input conditions because the axes were obtained from the manual instance annotations. With the upstream U-Net axes, the corresponding values were 0.438 and 0.452. The U-Net axis had recall/precision 0.782/0.753 on B58_100 and 0.729/0.887 on B58_110.

B58_100 was present in the segmenter’s training data; only B58_110 was verified training-clean. Moreover, each input mask defines a different visible common-fragment population, so the manual-derived and U-Net F_1 values are descriptive rather than paired. With two development fields, these observations show the practical dependence on upstream mask quality but do not estimate real-image generalisation.

4 Discussion

PLECTA addresses a narrow problem: assigning fragments in a binary filament-axis mask to overlapping projected instances. Its held-out common-fragment F_1 of 0.854 shows that explicit geometry can recover substantial identity without instance-labelled training data or SEM intensity. The mechanism is the combination of local and accumulated evidence. Exact matching at each junction selects mutually compatible continuations, gated gap matching reconnects separated arms, and repeated rounds replace short-arm direction estimates with tangent and

curvature estimates from the chains assembled so far. Shared crossing pixels remain in each applicable instance rather than being forced into one label.

The ablations support this interpretation but do not prove that every term is independently necessary. On the development set, removing the junction chord-turn term reduced F_1 by -0.1858, removing gap bridging by -0.1394, and disabling connector-based junction-cluster consolidation by -0.0634. Preventing chain-extended frames produced a smaller decrease of -0.0215. Thus the largest gains arise from chord-aware continuation, gap repair, and junction-topology preconditioning; iterative context provides an additional, more modest benefit. Because these tests reused development scenes, their values describe mechanism and configuration sensitivity, not independent effect estimates.

4.1 Density, masks, and optional rendering

Performance depends on the geometry preserved by the input mask. Mean held-out F_1 declined from 0.911 at 20% coverage to 0.775 at 60%, where crossings and locally plausible alternatives are more frequent. On 50 paired development scenes, the clean axis masks scored 0.0224 higher than degraded masks (95% CI 0.0113– 0.0333). PLECTA can bridge short omissions, but it cannot infer a filament absent from the mask or reliably reject a false axis introduced upstream. Supplying identical masks isolates grouping in an algorithmic comparison; it does not establish end-to-end SEM performance.

Image-informed ribbon rendering is downstream of grouping. On the development set it produced $F_1 = 0.860$, a change of -0.0068 relative to the core output, while yielding no disconnected instances and branch pixels in 0.005 of instances. This is a useful shape tradeoff rather than an improvement in identity. Silhouette absorption changes the instance partition and scored 0.857; it is therefore optional and should not be conflated with the fixed mask-only grouping result. Synthetic renderer widths were recovered accurately on purity-qualified development instances, but that validation concerns a known image model. Physical CNT diameter still requires independent, scale-matched references.

4.2 Evidence boundary

The principal limitation is distributional. The 128 held-out scenes use independent seeds but the same procedural generator family used for development. Their confidence intervals quantify sampling variation within that family, not transfer to another simulator, specimen, microscope, or segmentation pipeline. The ablation and rendering studies are development-only and must remain secondary to the fixed held-out result.

The real SEM study is deliberately a case study with two fields, not a population validation. Manual-derived masks are an oracle-like input condition: they test grouping when the axis evidence is curated, not automatic SEM-to-instance performance. B58_100 is training-contaminated for the U-Net condition, whereas B58_110 is documented as training-clean; only the latter can illustrate an unseen upstream mask, and one field cannot estimate generalization. Moreover, manual-derived and U-Net masks generate different fragment universes. Their F_1 values therefore describe performance under each input condition but are not a paired comparison of identical objects. More independent fields, sealed before configuration, and annotations from multiple readers are required.

Two further scope limits concern the output itself. First, PLECTA recovers projected two-dimensional identity. Shared pixels encode overlap but not which filament passes above another, so no three-dimensional topology follows. Second, the grouping metric evaluates ownership of centreline fragments and is not a validation of rendered width. Quantitative diameter claims require matched physical measurements across the relevant scale, contrast, and specimen ranges.

Pairwise fragment scores also emphasize instances containing many fragments; object recovery and ARI are retained as complementary views.

An independent procedural benchmark is a practical next test. The `merged_lines` mask generator in the CNT_SEM pipeline can export overlap-aware whole-bundle references and would add a different implementation and morphology model. It should supplement, not replace, the current generator: its field-of-view preset is presently calibrated to a narrow coverage regime, whereas the primary benchmark spans density. A multi-scene run should be made only after pooled calibration is fixed and its seeds are separated from any configuration work. No result from that generator is included in the present evidence. Broader validation should then extend to independent real masks and other filament domains while keeping input fragments and scoring rules fixed.

5 Conclusions

PLECTA reconstructs overlapping filament instances from a binary axis mask by alternating exact junction matching with gated gap matching as local arm geometry expands into chain-level evidence. The fixed method achieved common-fragment $F_1 = 0.854$ on 128 independently seeded held-out scenes from the same procedural generator, with performance declining as network density increased. Development ablations identify junction chord-turn evidence, gap bridging, and topology consolidation as the principal contributors. Optional image-informed rendering improves geometric presentation but remains separate from the grouping claim. Two real SEM fields demonstrate the workflow under manual-derived and U-Net masks, not population generalization or width accuracy. The method therefore provides an inspectable, deterministic approach for projected filament identity; broader claims require independent real fields, validated physical widths, and an external-generator benchmark with a fixed calibration and sealed seeds.

CRediT author contribution statement

[To do: Fill in each author’s CRediT roles before submission, checked against actual contributions; keep the author names and order as listed above.]

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data and code availability

The PLECTA source, complete fixed configuration, compact seed manifest, and aggregate machine-readable evidence are available at <https://github.com/Jorallan/PLECTA>. Generated scenes, detailed per-scene records, and development diagnostics are retained locally. [To do: Create a tagged archival release and add its DOI before submission; do not infer or preassign a DOI.]

Acknowledgements

The authors thank the maintainers of NumPy, SciPy, scikit-image, NetworkX, and Matplotlib. The upstream synthetic-data-driven mask workflow adapts the open CycleGAN/U-Net particle-

segmentation workflow of Rühle et al. [33]; it is not evaluated as the contribution of this paper.

References

- [1] Michael F. L. De Volder, Sameh H. Tawfick, Ray H. Baughman, and A. John Hart. Carbon nanotubes: Present and future commercial applications. *Science*, 339(6119):535–539, 2013. doi: 10.1126/science.1222453.
- [2] Jonathan N. Coleman, Umar Khan, Werner J. Blau, and Yurii K. Gun’ko. Small but strong: A review of the mechanical properties of carbon nanotube-polymer composites. *Carbon*, 44(9):1624–1652, 2006. doi: 10.1016/j.carbon.2006.02.038.
- [3] Alexey N. Volkov and Leonid V. Zhigilei. Structural stability of carbon nanotube films: The role of bending buckling. *ACS Nano*, 4(10):6187–6195, 2010. doi: 10.1021/nn1015902.
- [4] Igor Ostanin, Roberto Ballarini, David Potyondy, and Traian Dumitrică. A distinct element method for large scale simulations of carbon nanotube assemblies. *Journal of the Mechanics and Physics of Solids*, 61(3):762–782, 2013. doi: 10.1016/j.jmps.2012.10.016.
- [5] Bernard K. Wittmaack, Abu Horaira Banna, Alexey N. Volkov, and Leonid V. Zhigilei. Mesoscopic modeling of structural self-organization of carbon nanotubes into vertically aligned networks of nanotube bundles. *Carbon*, 130:69–86, 2018. doi: 10.1016/j.carbon.2017.12.078.
- [6] Liyu Dong, Haibin Hang, Jin Gyu Park, Washington Mio, and Richard Liang. Detecting carbon nanotube orientation with topological analysis of scanning electron micrographs. *Nanomaterials*, 12(8):1251, 2022. doi: 10.3390/nano12081251.
- [7] Tamjid Imtiaz, Jacques Doumani, Fuyang Tay, Natsumi Komatsu, Stephen Marcon, Motonori Nakamura, Saunab Ghosh, Andrey Baydin, Junichiro Kono, and Ahmed Zubair. Facile alignment estimation in carbon nanotube films using image processing. *Signal Processing*, 202:108751, 2023. doi: 10.1016/j.sigpro.2022.108751.
- [8] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-Net: Convolutional networks for biomedical image segmentation. In *Medical Image Computing and Computer-Assisted Intervention – MICCAI 2015*, volume 9351 of *Lecture Notes in Computer Science*, pages 234–241, 2015. doi: 10.1007/978-3-319-24574-4_28.
- [9] Fabian Isensee, Paul F. Jaeger, Simon A. A. Kohl, Jens Petersen, and Klaus H. Maier-Hein. nnU-Net: A self-configuring method for deep learning-based biomedical image segmentation. *Nature Methods*, 18:203–211, 2021. doi: 10.1038/s41592-020-01008-z.
- [10] Carsen Stringer, Tim Wang, Michalis Michaelos, and Marius Pachitariu. Cellpose: A generalist algorithm for cellular segmentation. *Nature Methods*, 18:100–106, 2021. doi: 10.1038/s41592-020-01018-x.
- [11] Uwe Schmidt, Martin Weigert, Coleman Broaddus, and Gene Myers. Cell detection with star-convex polygons. In *Medical Image Computing and Computer Assisted Intervention – MICCAI 2018*, *Lecture Notes in Computer Science*, pages 265–273, 2018. doi: 10.1007/978-3-030-00934-2_30.
- [12] Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, Tete Xiao, Spencer Whitehead, Alexander C. Berg, Wan-Yen Lo, Piotr Dollár, and Ross Girshick. Segment anything. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 3992–4003, 2023. doi: 10.1109/ICCV51070.2023.00371.

- [13] Carsten Steger. An unbiased detector of curvilinear structures. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 20(2):113–125, 1998. doi: 10.1109/34.659930.
- [14] Alejandro F. Frangi, Wiro J. Niessen, Koen L. Vincken, and Max A. Viergever. Multiscale vessel enhancement filtering. In *Medical Image Computing and Computer-Assisted Intervention – MICCAI’98*, volume 1496 of *Lecture Notes in Computer Science*, pages 130–137, 1998. doi: 10.1007/BFb0056195.
- [15] Yoshinobu Sato, Shin Nakajima, Nobuyuki Shiraga, Hideki Atsumi, Shigeyuki Yoshida, Thomas Koller, Guido Gerig, and Ron Kikinis. Three-dimensional multi-scale line filter for segmentation and visualization of curvilinear structures in medical images. *Medical Image Analysis*, 2(2):143–168, 1998. doi: 10.1016/S1361-8415(98)80009-1.
- [16] Andrew M. Stein, David A. Vader, Louise M. Jawerth, David A. Weitz, and Leonard M. Sander. An algorithm for extracting the network geometry of three-dimensional collagen gels. *Journal of Microscopy*, 232(3):463–475, 2008. doi: 10.1111/j.1365-2818.2008.02141.x.
- [17] Jeremy S. Bredfeldt, Yuming Liu, Carolyn A. Pehlke, Matthew W. Conklin, Joseph M. Szulczewski, David R. Inman, Patricia J. Keely, Robert D. Nowak, Thomas R. Mackie, and Kevin W. Eliceiri. Computational segmentation of collagen fibers from second-harmonic generation images of breast cancer. *Journal of Biomedical Optics*, 19(1):016007, 2014. doi: 10.1117/1.JBO.19.1.016007.
- [18] Ting Xu, Dimitrios Vavylonis, Feng-Ching Tsai, Gijsje H. Koenderink, Wei Nie, Eddy Yusuf, I-Ju Lee, Jian-Qiu Wu, and Xiaolei Huang. SOAX: A software for quantification of 3D biopolymer networks. *Scientific Reports*, 5:9081, 2015. doi: 10.1038/srep09081.
- [19] Daniel A. D. Flormann, Moritz Schu, Emmanuel Terriac, Divyendu Thalla, Lucina Kainka, Marcus Koch, Annica K. B. Gad, and Franziska Lautenschläger. A novel universal algorithm for filament network tracing and cytoskeleton analysis. *The FASEB Journal*, 35(5):e21582, 2021. doi: 10.1096/fj.202100048R.
- [20] Nathan A. Hotaling, Kapil Bharti, Haydn Kriel, and Carl G. Simon. DiameterJ: A validated open source nanofiber diameter measurement tool. *Biomaterials*, 61:327–338, 2015. doi: 10.1016/j.biomaterials.2015.05.015.
- [21] Dorian Depriester, Sabine Rolland du Roscoat, Laurent Orgéas, Christian Geindreau, Benjamin Levrard, and Florian Brémont. Individual fibre separation in 3D fibrous materials imaged by X-ray tomography. *Journal of Microscopy*, 286(3):220–239, 2022. doi: 10.1111/jmi.13096.
- [22] Saurav Basu, Chi Liu, and Gustavo Kunde Rohde. Extraction of individual filaments from 2D confocal microscopy images of flat cells. *IEEE/ACM Transactions on Computational Biology and Bioinformatics*, 12(3):632–643, 2015. doi: 10.1109/TCBB.2014.2372783.
- [23] David Breuer and Zoran Nikoloski. DeFiNe: An optimisation-based method for robust disentangling of filamentous networks. *Scientific Reports*, 5:18267, 2015. doi: 10.1038/srep18267.
- [24] Jaydeep De, Li Cheng, Xiaowei Zhang, Feng Lin, Huiqi Li, Kok Haur Ong, Weimiao Yu, Yuanhong Yu, and Sohail Ahmed. A graph-theoretical approach for tracing filamentary structures in neuronal and retinal images. *IEEE Transactions on Medical Imaging*, 35(1): 257–272, 2016. doi: 10.1109/TMI.2015.2465962.
- [25] Zhen Zhang, Yukako Nishimura, and Pakorn Kanchanawong. Extracting microtubule networks from superresolution single-molecule localization microscopy data. *Molecular Biology of the Cell*, 28(2):333–345, 2017. doi: 10.1091/mbc.E16-06-0421.

- [26] Clément Ployat, Yosra Mehrjoo, Renaud Duval, Marie Carole Boucher, Santiago Costantino, and Hugo Wurtele. DNAi: An open-source AI tool for unbiased DNA fiber analysis. *Nucleic Acids Research*, 54(7):gkag335, 2026. doi: 10.1093/nar/gkag335.
- [27] Yi Liu, Abhishek Kolagunda, Wayne Treible, Alex Nedo, Jeffrey Caplan, and Chandra Kambhamettu. Intersection to overpass: Instance segmentation on filamentous structures with an orientation-aware neural network and terminus pairing algorithm. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 125–133, 2019. doi: 10.1109/CVPRW.2019.00021.
- [28] Viktor Wegmayr, Aytunc Sahin, Björn Sæmundsson, and Joachim M. Buhmann. Instance segmentation for the quantification of microplastic fiber images. In *2020 IEEE Winter Conference on Applications of Computer Vision (WACV)*, pages 2199–2206, 2020. doi: 10.1109/WACV45572.2020.9093352.
- [29] Max Frei and Frank Einar Kruis. FibeR-CNN: Expanding Mask R-CNN to improve image-based fiber analysis. *Powder Technology*, 377:974–991, 2021. doi: 10.1016/j.powtec.2020.08.034. Published online 2020.
- [30] Yi Liu, Su Peng, Jeffrey Caplan, and Chandra Kambhamettu. Pick and trace: Instance segmentation for filamentous objects with a recurrent neural network. In *Medical Image Computing and Computer Assisted Intervention – MICCAI 2023*, Lecture Notes in Computer Science, pages 635–645, 2023. doi: 10.1007/978-3-031-43993-3_61.
- [31] Isabella Østerlund, Arne Neumann, Zhiming Ma, Yansong Miao, Staffan Persson, and Zoran Nikoloski. GraFT: A robust network-based spatiotemporal analysis of filamentous structures. *Science Advances*, 12(13):eadz4132, 2026. doi: 10.1126/sciadv.adz4132.
- [32] Torben Peters, John Schumann, Kerstin Kämpf, and Asmus Meyer-Plath. AI-assisted recognition of fibre instances in scanning electron microscope images for aerosol characterization and exposure assessment. *Powder Technology*, 469:121909, 2026. doi: 10.1016/j.powtec.2025.121909.
- [33] Bastian Rühle, Julian Frederic Krumrey, and Vasile-Dan Hodoroaba. Workflow towards automated segmentation of agglomerated, non-spherical particles from electron microscopy images using artificial neural networks. *Scientific Reports*, 11(1):4942, 2021. doi: 10.1038/s41598-021-84287-6.
- [34] Oday Allan. CNT_SEM: CycleGAN synthesis and U-Net segmentation workflows for carbon-nanotube SEM data. GitHub repository, 2026. URL https://github.com/Jorallan/CNT_SEM.
- [35] William M. Rand. Objective criteria for the evaluation of clustering methods. *Journal of the American Statistical Association*, 66(336):846–850, 1971. doi: 10.1080/01621459.1971.10482356.
- [36] Ignacio Arganda-Carreras, Srinivas C. Turaga, Daniel R. Berger, Dan Cireşan, Alessandro Giusti, Luca M. Gambardella, Jürgen Schmidhuber, Dmitry Laptev, Sarvesh Dwivedi, Joachim M. Buhmann, Ting Liu, Mojtaba Seyedhosseini, Tolga Tasdizen, Lee Kamensky, Radim Burget, Václav Uher, Xiao Tan, Changming Sun, Tuan D. Pham, Erhan Bas, Mustafa G. Uzunbas, Albert Cardona, Johannes Schindelin, and H. Sebastian Seung. Crowdsourcing the creation of image segmentation algorithms for connectomics. *Frontiers in Neuroanatomy*, 9:142, 2015. doi: 10.3389/fnana.2015.00142.
- [37] Lawrence Hubert and Phipps Arabie. Comparing partitions. *Journal of Classification*, 2(1): 193–218, 1985. doi: 10.1007/BF01908075.

- [38] Marina Meilă. Comparing clusterings—an information based distance. *Journal of Multivariate Analysis*, 98(5):873–895, 2007. doi: 10.1016/j.jmva.2006.11.013.