

PLECTA: Overlap-aware instance reconstruction of dense filament networks from binary masks

Oday Allan^a Stefan Luding^a Igor Ostanin^{a,*}

^a Multi-Scale Mechanics (MSM), Faculty of Engineering Technology, University of Twente, Enschede, The Netherlands

[To do: Add ORCIDs.]

Abstract

A dense network of filaments defeats instance segmentation in two opposite ways at once. Where filaments cross, separate objects fuse into one connected component; where the segmentation fades, one object breaks into several. Both faults look identical in a binary mask, and repairing either requires the same judgement: which arm continues into which.

We present PLECTA, which makes that judgement from a binary filament-axis mask and nothing else. It cuts the skeleton into arms and junctions and solves every junction, of any degree, as a single exact maximum-weight *partial* matching in which leaving an arm unmatched carries an explicit price. A filament is therefore allowed to end, instead of having a continuation invented for it. Crossing pixels may belong to more than one instance, so overlap is represented rather than partitioned away. The evidence is purely directional, tangent and curvature read over tens of pixels either side of the discontinuity, which makes the working condition plain: the method applies where direction survives a crossing, as it does for nanotube bundles, cytoskeletal filaments and synthetic fibres, and not for freely flexible chains.

On 128 held-out synthetic scenes PLECTA reaches mean pairwise $F_1 = 0.854$ [0.845, 0.864], falling from 0.911 to 0.775 as coverage rises from 20 to 60 %, with no failed scenes. It leads two released implementations, run from their own sources on identical masks, at every density. Because such a lead could be an artefact of our own generator, the comparison was repeated on a competitor's: GraFT's, vendored unmodified at its published settings over 40 scenes. The ordering held at every density stratum, at F_1 0.916 against 0.475 and 0.349.

An optional stage then reads the greyscale to decide which filament passes in front at each crossing, and stacks the network into layers, information otherwise obtained by destructive sectioning. It is correct on 0.861–0.908 of the crossings it commits to, against a chance level of 0.5, and abstains where the image does not separate the two. The method uses no random number generator and its configuration is published, so a mask yields the same instances on every run. We demonstrate the whole workflow on 3 annotated carbon-nanotube SEM fields, and state plainly what that evidence does not support.

Keywords: filament networks; instance segmentation; graph matching; geometric reconstruction; carbon nanotubes; scanning electron microscopy.

1 Introduction

Many materials are networks of long, thin filaments. Sheets of carbon nanotubes, the cytoskeleton inside a living cell, paper, felt and synthetic fibre mats all consist of slender objects laid over one another in a tangle. What such a material does (how stiff it is, how well it conducts, how it tears) depends less on the filaments themselves than on the geometry of the network they form: how long each one is, how thick, which way it runs, how sharply it bends, and which others it touches [1–5].

*Corresponding author: i.ostanin@utwente.nl

A micrograph shows that network directly. Turning the picture into those numbers is the hard part. Routine image analysis measures the network in bulk (what fraction of the frame is covered, which way the texture points [6, 7]) and stops there. To measure a distribution of filament lengths, or to ask how often one filament crosses another, something more is needed: every filament must be recovered as a separate object, addressable from one end to the other.

That is an instance-segmentation problem, and it has an awkward geometry. The objects are long, thin and open rather than compact and closed; they touch, cross and hide one another; and what tells two of them apart is not shape or size but which way each was heading. The consequence is a pair of failures that pull in opposite directions. Where two filaments cross, their pixels join into a single connected blob, and a method that trusts connectivity reports one object where there are two. Where the image fades or the segmentation drops out, one filament arrives in several disconnected pieces, and the same method reports several objects where there is one. In a binary mask the two look identical, and both are repaired by answering one question: at this junction, which arm continues into which?

Existing tools answer part of it. Semantic segmentation is mature, since U-Net and its self-configuring descendants reliably separate filament foreground from background [8, 9], but a foreground mask carries no identity, so it inherits both failures rather than resolving them. General-purpose instance methods assume the objects are compact: bounded blobs, centre-directed flows, star-convex outlines, none of which describes a long open curve that may overlap another [10–12]. Ridge and vesselness filters sharpen curvilinear signal but assign nothing across a junction [13–15], and the established fibre-analysis packages return network statistics rather than resolved objects [16–20]. This paper addresses the step those leave open: given a binary filament-axis mask, recover the individual filaments through the crossings and the gaps.

Recovering filaments through crossings is itself well-trodden ground, and the design space is thoroughly explored. Basu et al. build a spanning tree over centreline particles, delete every bifurcation, and recombine the segments with an integer program priced by spline curvature, re-solving until the segment set stops changing [21]. DeFiNe consumes a pre-extracted weighted graph and covers it with paths of minimum roughness, optionally letting several paths share an edge so that overlaps are representable [22]. SIFNE excises each crossing, estimates tip directions, and pairs termini greedily inside an angular gate [23]. Others partition a directed graph [24], group skeleton branches by lowest mutual angle and learn pixel embeddings for tangled microplastic fibres [25], enumerate the pairings at junctions of degree three or four from endpoint angles [26], or trace a Frangi-filtered skeleton by constrained search and a greedy longest-first path cover [27]. Learned alternatives pair termini with an orientation network, detect fibres with a region-based detector, or trace them recurrently [28–30].

Two of these are compared against directly here, so what each was built for matters, because neither was built for this. DNAi is an end-to-end tool for DNA fibre spreading assays, where the quantity wanted is a length ratio between two labelled tracts, and it matches trained human annotators on that task [26]. GraFT traces and tracks the plant actin cytoskeleton through confocal time series, a temporal dimension outside this paper’s scope [27]. Both are strong in their domains, and both make choices those domains license: DNAi commits to a pairing at every junction of degree two to four and attempts none above that, while GraFT’s greedy path cover cannot revise an early commitment. Such choices cost little where filaments are sparse and well separated, which is the regime each was validated in, and more where filaments are dense and crossings frequent. Section 3.3 runs both from their own released sources, on identical masks, through one scorer.

Skeleton graphs, junction excision, continuation angles, distance gates, combinatorial optimisation, iterative refitting and overlapping output are therefore all prior art, and none is claimed here. What is new is narrower, and it is a combination: from a binary mask alone, PLECTA solves each junction of *any* degree as one exact maximum-weight *partial* matching in which leaving an arm unmatched is an explicit option at a fixed price rather than a parity leftover,

alternates that with globally gated gap matching, re-fits its tangent and curvature frames along the chains as they assemble, and emits overlapping instances. Each element exists somewhere above; no method there combines them. The globally optimising methods consume intensity evidence or a pre-extracted graph rather than a mask, and recombine chains of at most three segments [21, 22]; the mask-reading methods pair endpoints greedily [23, 25, 27] or stop at degree four [26]. Degree-general exact matching with a priced refusal is what lets a filament genuinely end rather than have a continuation invented for it. The closest single prior method is that of Liu et al., which likewise reads a mask, reconnects across gaps and permits overlapping output, but pairs termini greedily and depends on a trained orientation network [28].

That claim carries a condition, worth stating as scope rather than leaving to be inferred. The only evidence PLECTA has at a crossing is directional: the angle between two outward tangents, the turns onto the chord between them, and the continuity of signed curvature, each read over tens of pixels of arclength. The working condition is therefore a property of the image: direction must persist across the junction at that scale, or the continuation is not decidable from geometry at all. Bending stiffness is what most directly produces such persistence, which is why the semiflexible regime is the method’s home ground: nanotube bundles sit there, and so do the actin, collagen and synthetic-fibre networks the methods above were built for. Stiffness contributes to the condition rather than uniquely causing it, since tension, adhesion, confinement or bundling can hold a direction just as straight. Where the condition fails, as for a freely flexible chain whose tangent decorrelates within the junction, no continuation is more plausible than another, and no amount of matching recovers what the image no longer distinguishes.

We introduce PLECTA (Pairwise Linking with Evidence from Chain Tangents and Assignment; also the Latin *plecta*, a braid). It decomposes a skeleton into arms, endpoint stubs and merged junction nodes, alternates exact per-node continuation matching with globally gated gap matching while tangent and curvature estimates expand from local arms to assembled chains, and returns projected 2-D instance layers in which crossing pixels may be shared. Every parameter is fixed in a published configuration and the implementation draws on no random number generator, so a given mask yields the same instances on every run; that is the only sense in which it is called deterministic here.

Reaching real micrographs takes one further component, and the workflow is deliberately divided in two. An upstream route turns a raw SEM micrograph into a binary filament-axis mask by a synthetic-data-driven CycleGAN/U-Net whose training pairs are synthesised rather than annotated [31, 32]; PLECTA is everything downstream of that mask. The division is what makes the grouping measurable on its own, against references whose identities are known exactly, and it keeps width rendering from moving an identity score. The closest complete workflow is that of Peters et al., who join a U-Net-like segmentation trained on over a thousand annotated SEM micrographs to algorithms that trace fibres through intersections, and validate against human evaluators on fibre number, length and width [33]. That work and this one divide the task at the same joint and differ in what the instance step is asked to prove: their tracer stands in for human counting and measurement, while PLECTA is isolated behind a binary-mask contract and scored on identity through crossings against overlap-aware references in which every pairing is known. We evaluate PLECTA on 128 unseen scenes from the development generator and on 50 paired clean and degraded scenes, and demonstrate the two stages together on 3 annotated CNT-sheet SEM fields, 2 of them held out from the upstream segmenter.

Beyond the plane, a crossing carries a second question the grouping does not answer: which filament passes in front. For nanostructured networks that is normally obtained destructively, by focused-ion-beam sectioning and tomographic reconstruction [34], because a micrograph’s intensities record surface topography rather than height. PLECTA answers it optionally and non-destructively from the greyscale already in hand. The machinery is deliberately conventional: monocular depth ordering has long derived pairwise depth from local occlusion evidence, assembled those relations into a depth-order graph and corrected the contradictions




Figure 1. Where PLECTA sits. The method reads one binary axis mask and nothing else; every result in this paper measures the dashed green box. Steps 3 and 4 are the only decisions it makes and are drawn heavier for it, the rest is bookkeeping. The optional depth stage adds a third dimension without revising any identity. Colours: green, a PLECTA geometric step; orange, a dependency on the grayscale image; dark, a data object.

among them [35, 36], and single-view hair modelling resolves the layering of intersecting strands in much the same spirit [37]. What differs is the evidence available. Those methods key on T -junctions, where one contour terminates against another; two crossing filaments produce no such junction, since both contours continue through and the mask is a single blob. The cue used here is instead the continuity of a filament’s own appearance across the crossing, where the upper filament’s brightness and edge sharpness run through uninterrupted, while the lower filament’s are overwritten, and the stage abstains where that evidence does not separate the two. To our knowledge this is the first depth decomposition of a filament network from a single micrograph; it is reported as a demonstration with a measured accuracy, not as a validated three-dimensional reconstruction.

2 Materials and methods

2.1 Scope and workflow

The formal input is a binary filament-axis mask $M \in \{0, 1\}^{H \times W}$, from which PLECTA re-constructs an overlap-aware collection $\{I_1, \dots, I_N\}$ with $I_n \subseteq M$ and a crossing pixel possibly belonging to several instances. The evaluated grouping stage uses only M ; a grayscale SEM image is optional downstream input for width and brightness measurement and cannot revise the fixed grouping.

Figure 1 places the method. Its five geometric steps read mask geometry alone, and the two matching steps are re-made over eight rounds against frames re-fitted along the chains assembled so far. Where the mask itself comes from is a separate question, taken up in Section 3.7.

What makes the task hard is a single crossing: the mask holds one connected blob from which several arms leave, and nothing local states which arm continues into which. The output is a set of layers rather than a partition, so a pixel inside that blob is counted in both instances that claim it.

2.2 Synthetic data and evaluation split

Procedural scenes preserve the ordered centreline $\mathcal{C}_k$ and rendered support Ω_k of every filament k ; supports remain separate, so several filaments may own the same crossing pixel. Besides the exact overlap-aware reference, the generator provides an undegraded one-pixel axis mask, a degraded thin mask containing small gaps and irregularities, a combined thick foreground, and an SEM-like image. The degraded one-pixel axis is the input evaluated throughout. Areal coverage is the union area of the thick supports divided by image area, not the occupancy of the thin input mask, and Section 3.2 shows it is a proxy for difficulty rather than a cause of it.

Development used 84 scenes spanning 20–60% coverage. The held-out set comprised 128 different seeds, 32 scenes at each of 20, 30, 40, and 60% coverage: an independent sample from the same generator, not an independent distribution. A separate 125-scene locked set remains unevaluated. The compact seed manifest records every seed block and an executable verifier checks counts, disjointness, and available scene metadata.

A separate controlled paired set of 50 scenes, ten at each of 20, 30, 40, 50, and 60% coverage, was scored once from each geometry’s clean axis and once from its degraded mask, after the configuration was fixed.

2.3 CNT SEM data and mask sources

3 manually annotated CNT-sheet SEM fields were available for exploratory analysis, each 1536×1024 px: B58_100 (307 visible CNT bundles), B58_110 (430) and B58_300 (325). The annotated instance is the visible bundle, not the individual nanotube, which SEM does not resolve at this magnification. The nominal horizontal field width is $1.66 \mu\text{m}$ across the long, 1536-px axis of the stored images, which is 1.08 nm px^{-1} . Each annotation is one human interpretation of projected, locally ambiguous crossings, and B58_100 and B58_110 informed development.

The reference annotations were drawn with an in-house interactive instance-segmentation tool for greyscale SEM of thin, crossing filaments, whose data model is per-instance masks rather than a flat label image, so a pixel may belong to several instances at once; that is what makes the reference overlap-aware and comparable with PLECTA’s output at all. Its primary artefact, a compressed sparse per-instance pixel set, is the identical file the evaluation loads as its reference, so provenance is end to end. The annotator paints each visible bundle as its own instance, and the tool assists without deciding: strokes may snap onto the local intensity ridge, with polarity calibrated from the existing annotation, and may be re-rendered through a smoothing-spline fit whose centreline is shifted onto the measured intensity crest only where that crest is consistent and by at most about 8 px; widths come from perpendicular full-width-at-half-maximum profiles referenced to the local minima either side of the peak and reported as twice the nearer half-width, so a crossing filament fusing into the profile does not bias the measurement wide. Candidate gap bridges, collinear endpoint pairs of different instances, are suggested but never applied automatically; each is accepted or rejected by the annotator, so identity through every crossing is a human decision, machine-assisted but not machine-made. The tool keeps an audit log of every mutating action alongside the saved annotation, and an advisory per-instance plausibility score, computed from profile support, peak contrast and width plausibility, was used as quality control while annotating; the score is advisory only and is not part of the saved reference.

The protocol’s limits are stated rather than implied. The primary annotator is an author of this paper and also configured PLECTA, and there was no blinding, so differences smaller than plausible annotation variation should not be read as real, and the annotations are manual reference annotations, not ground truth. A second reader has since annotated 3 of these fields independently, and Section 4.1 reports what that measurement shows: the two readers do not implement one instance definition, so every real-data score in this paper is reported against a named reader and the two are never averaged. One assumption is also shared across the divide:

the annotator’s width measurement and PLECTA’s optional width layer of Section 2.5 both read perpendicular FWHM profiles across a centreline. That does not touch the grouping scores, which are computed on centreline fragments, but it does mean the two are not independent about width; nothing in this paper compares the two widths, and no such comparison should be made without restating this.

Two inputs were derived per field. The first is the skeleton of the manual instance annotation, an oracle-like condition that asks what the grouping can do when the axis evidence is curated. The second is an automatic mask from the upstream CycleGAN/nnU-Net segmentation of Section 3.7 [9, 31, 32], which asks what an automatic mask allows it to do. That segmentation is a separate contribution with its own evaluation; all that matters here is that its output is a binary mask produced without seeing any annotation.

Which fields that model had seen is a property of the (model, field) pair and was verified by image content rather than by filename. B58_110 and B58_300 are absent from the training set of the run used here; B58_100 is present in every available run and its U-Net condition is therefore an upper bound throughout. Section 3.7 measures what that contamination is worth rather than leaving it as a caveat. Mask quality was compared after skeletonisation with the symmetric 2-px tolerance used throughout, which avoids rewarding mask thickness, and reconstruction scores were computed within each mask’s own visible common-fragment population.

[Needed before submission: Author metadata: SEM instrument and detector, accelerating voltage, working distance, specimen preparation, and CNT material identity.]

2.4 PLECTA reconstruction

2.4.1 Skeleton graph and local frames

PLECTA skeletonises the mask and prunes the short spurs skeletonisation leaves behind. Skeleton pixels where three or more paths meet form junction clusters, and clusters close enough together are merged into a single *node*: a crossing seen at this resolution is one blob, not a point, and neighbouring crossings often touch. What remains between nodes is an *arm*, an unambiguous run of filament, carrying a *stub* at each end.

Each stub is given a frame: its tip, the direction the filament is heading as it leaves, and how it is curving. Both are fitted over a short run of arclength back from the tip, and a stub too short to yield a reliable direction is marked as such rather than guessed at. That support widens once arms are linked into chains, so a stub with no direction of its own borrows one from the chain it has joined, which is why the method iterates rather than deciding once.

Linking two stubs is then priced by asking three geometric questions (Figure 2). Do the two directions oppose one another, as they would if one filament simply continued into the other? Does each direction point at the other’s tip, rather than running parallel past it? And does the curvature continue, or would the join require a kink? For a gap the distance to be spanned is charged as well; at a junction there is no gap to pay for. The cost of linking stubs a and b is therefore

$$C_{ab} = w_d \theta + w_t q(d) (\varphi_a + \varphi_b) + w_\ell \frac{d}{\ell_0} + 30 w_\kappa |\kappa_a + \kappa_b|, \quad q(d) = \min\left(1, \frac{d}{d_0}\right), \quad (1)$$

with d the tip separation, θ the reversal angle between the two outward tangents, and φ_a, φ_b the turns from each tangent onto the tip-to-tip chord, measured at their own vertices; junction links set $w_\ell = 0$.

Three properties of that expression are deliberate. The reversal angle ignores the chord, which keeps it meaningful when the tips are three pixels apart and the chord direction is quantisation noise. The chord turns do see lateral offset, so two parallel arms that never meet are penalised. And because it is the chord terms that degrade over short separations, the fade $q(d)$ withdraws




Figure 2. The stub frame and the terms of the pairing cost. (a) The least-squares arclength fit over window w , read outward from the tip at $s = 0$, giving outward tangent $\mathbf{t}$ and signed curvature κ ; the fit is quadratic when the span is at least 18px and linear otherwise. (b) $\theta = \arccos(\mathbf{t}_a \cdot -\mathbf{t}_b)$, in the two steps its definition takes: the two outward tangents where they lie, then the same pair with $\mathbf{t}_b$ reversed and both read from one origin, which is where the angle is measured. A perfect continuation makes the two tangents antiparallel, so $\theta = 0$. (c) The chord turns φ_a, φ_b onto the chord $\mathbf{u}$ joining the tips. (d) The fade $q(d)$ against tip separation; triangles mark the six real separations at the crossing of Figure 3.

them as the tips approach rather than letting an unstable angle dominate. The weights, the scales ℓ_0 and d_0 , and the constant placing curvature on the scale of the angular terms are all fixed in the released configuration and are listed in Supplementary Section S1.

2.4.2 Candidate gates and exact matching

Leaving a stub unmatched is priced rather than forbidden, and that price does double duty: it also fixes what may be offered at all, since an edge costing more than the two prices it would save could never be chosen. Exact maximum-weight partial matching then picks, over disjoint pairs,

$$\max \sum_{(i,j) \in \mathcal{M}} (p_i + p_j - C_{ij}), \quad \text{admissible iff } C_{ij} < p_i + p_j, \quad (2)$$

which is the same thing as minimising the total link cost plus a price p_i for every stub left free. Junction and gap links carry different prices, and a gap link must additionally clear limits on how far it may reach and how far either direction may turn, the tighter set, because a gap asserts a continuation across evidence that is *missing* rather than merely ambiguous. Supplementary Section S1 lists the prices, the limits, and the schedule below.

The method does not decide once. It runs the two matchings over several rounds, re-fitting every frame along the chains assembled so far before each. Early rounds are deliberately conservative: prices are scaled down, which admits fewer edges, and the gap angle limits are narrowed, so only pairings that survive the stricter test are committed while the frames are least reliable. The constraints relax towards their configured values as the chains sharpen the frames, and only the final round runs at the configured parameters. That round is the one returned.

Instances are assumed to be open projected curves, and a guard enforces it: any matched component whose links close a ring has its longest link removed. On the 84 development scenes the guard never fired, with zero links removed, so it protects against a configuration the data did not produce rather than shaping the result. Closed or looping filaments are outside the scope of this paper.




Figure 3. Every decision the linker can make, on the nodes it made them at. Every number shown is measured, and every arm is the skeleton it was measured from, drawn through a short moving average that removes the raster staircase and moves no point by a pixel; the two nodes are rigidly rotated to fit their panels, so every angle shown is the measured one. (a) A four-arm crossing (development scene cov20/synth_0001, node 9) before and after the decision: arms named *a* to *d* anticlockwise by outward tangent, the junction cluster shaded, and the two instances it becomes, which share seven pixels. (b) Every configuration the admissibility gate allows at that node, each scored as a whole. Under each is the sum of its edge costs and the price $p_i = 0.62$ of every arm it leaves free, a ring marking such an arm; the boxed configuration is the one returned. The two candidate pairs missing here cost 3.15 and 3.33, above $p_i + p_j = 1.24$, so pairing them would cost more than leaving both free. (c) The same at a three-arm node (cov40/synth_0001), where an admissible edge costing 1.01 is declined: taking it totals 1.63, against 0.77 for pairing the two arms that continue into one another and paying to leave the third free, and 1.86 for leaving all three free. (d) The third decision, on the same scene: two arm tips 25.8 px apart joined across a break in the mask, with the other mask pieces of the crop in grey. A gap link must clear all four gates at once.

Every stub is therefore either paired or pays the unmatched price, and there is no third possibility. That priced option is what lets a filament genuinely end, and what stops a T-junction from inventing a continuation for the arm that has no partner.

Solving a junction as one matching rather than one edge at a time is the distinction claimed in Section 1, and Figure 3 shows what it buys on real nodes, alongside the other two decisions the linker makes. A locally attractive edge can be declined, as at the three-arm node of Figure 3c. Joint and sequential solutions coincide at plain crossings and diverge only as nodes grow dense: they never disagree at nodes of degree two to four, disagree at roughly a fifth of nodes with five to eight stubs, and at two thirds of nodes with nine or more, where the joint solution is the cheaper one; Supplementary Section S1 gives the counts. That comparison holds one node's costs fixed. Run inside the rounds, where each matching sets the frames the next is scored against, a cheap early choice propagates and the difference is no longer confined to dense nodes (Section 3.3).

2.4.3 Instance output

Connected components of accepted arm links define instances. PLECTA paints every arm pixel and every junction cluster touched by a chain into that instance, so intersecting instances share crossing pixels. An isolated one-arm component shorter than 6 px is omitted. Accepted gap links determine identity but are not painted into the fixed centreline output. All effective parameters, including values previously inherited as code defaults, are materialised in the fixed JSON configuration.

2.5 Image-informed width reconstruction, optional

The grouping above reads the mask alone and is what this paper evaluates. Where a grayscale image of the same field exists, two further stages can use it to give each instance a width, and both are off by default and governed by limits recorded in the fixed configuration. Neither is specific to SEM: they take whatever image accompanies the mask, and SEM is simply the case demonstrated here.

Perpendicular intensity profiles, sampled along each ordered chain away from junctions, attach an apparent width, a brightness, an uncertainty and a quality flag to an existing instance without altering its arm membership. Ribbon rendering then interpolates accepted gaps, resamples and smooths the centreline, and draws the chain at its fitted width, excluding junction blobs. Silhouette absorption, which would union rendered instances, has threshold zero by default; neither it nor any post-hoc continuation merge is part of the evaluated algorithm or contributed to the held-out result. Because the grouping is fixed before either stage runs, no width measurement can move a pixel from one instance to another, and no identity score in this paper depends on one.

2.6 Depth decomposition, optional

A projected crossing has an owner in each of two instances but no order between them. This stage supplies one. It runs after the grouping is fixed, reads the grayscale and the instances, and changes no instance membership.

Evidence at one crossing. Where two instance centrelines intersect, a short arc of each is taken through the crossing (the shared *core*, sized from the two radii), together with the *flanks* on either side of it, kept clear of the other filament’s stroke so that they measure one filament and not the union. Two channels are read on both: median intensity, and median gradient magnitude on the filament’s two edges. The physical argument is occlusion. Whichever filament lies on top presents the same surface through the crossing as it does beside it, so its core matches its own flanks in brightness and its edges stay crisp; the lower filament is overwritten there, so both channels break. The filament whose flank appearance the core matches better is called upper, and the normalised difference between the two channels passes through a logistic to give p . Where the two filaments genuinely look alike the difference is near zero, p sits near 0.5, and the crossing *abstains* rather than being decided by noise.

One order for the whole field. Local decisions need not be mutually consistent: they can assert that A is above B , B above C and C above A , which no assignment of heights satisfies. Each decided crossing therefore contributes a directed edge weighted by its own confidence, and within each connected component of the crossing graph the orientation that maximises total satisfied weight is chosen, equivalently a minimum-weight feedback arc set. Relations that cannot be satisfied are flipped rather than discarded, and both the flip and the original local probability are recorded, so a reader can see how far the global correction moved. Depth propagates only along crossings, so two instances in different components are related by no evidence at all and their order is a stated tie-break, not a measurement. Supplementary Section S2 gives the sampling geometry, the weighting and the solver.

Layers and height. The corrected relations form a directed acyclic graph whose longest path fixes the smallest number of layers consistent with them; filaments that never conflict share a layer, so the layer count follows from the ordering rather than from the number of instances. Each instance is then placed at a height that respects its order and keeps filaments from interpenetrating at their measured radii, with the stack compacted so that the film is as thin as those constraints allow.

Without a greyscale image every crossing abstains, the ordering has nothing to satisfy, and the output degenerates to the tie-break: the stage requires image evidence and reports how much of it was usable. Section 3.8 evaluates it, and Supplementary Section S5 defines the measures used, which are disjoint from those above because none of them can see depth.

2.7 Metrics and statistical analysis

The primary endpoint is common-fragment pairwise F_1 . The input mask is cut into fragments, each one unambiguous arm between crossings, by a rule that reads neither the reference nor any prediction, so every method compared is cut into the same pieces; each fragment then takes a reference and a predicted identity by one rule applied to both sides alike, and the score counts the pairs of fragments the prediction groups as the reference does. Supplementary Section S3 gives the construction and the assignment rule in full. Two points bear on the interpretation. The scorer shares two constants with the method: both prune skeleton spurs of at most 3 px, and the scorer’s 6 px minimum fragment length coincides with the length below which PLECTA omits an isolated component, so endpoint and method rest on a shared preprocessing assumption rather than independent ones, and a prespecified one-at-a-time sweep of the four constants is reported in Supplementary Section S3. And because junction pixels are deleted before fragments are labelled, the measure cannot see whether an emitted crossing pixel is given one owner or two, only whether the arms meeting there are grouped as the reference groups them; Section 3.4 separates the two.

Several measures accompany it, and the reason for carrying more than one is a single distinction: *pairwise scores relationships between pieces, detection scores the existence of objects*. A filament delivered in n pieces contributes $\binom{n}{2}$ pairs but one object, so pairwise F_1 is quadratic in fragmentation and a broken mask is charged many times over for one missed join, while object detection is blind to an arm swap at a crossing, scoring an inverted reconstruction 1.000 where pairwise scores 0.000; and neither states how many crossings were solved. So the adjusted Rand index and variation of information [38, 39] are quoted beside the pairwise score, detection F_1 beside them, join F_1 where two mask sources fragment differently, and reference-instance recovery to weight every filament equally rather than by its pair count. Crossing Fidelity, the fraction of reference crossings at which the predicted grouping of the incident arms matches the reference partition exactly, says how many crossings were solved, and is an in-house diagnostic rather than an established metric, in the tradition of measures that score correctness at branch points rather than per pixel [40, 41]; it is always printed beside its own all-unpaired baseline, the fidelity of the deterministic do-nothing policy that leaves every arm unpaired. The input masks themselves, as distinct from the reconstructions, are characterised by the tolerant completeness/correctness/quality triple standard for thin structures. Supplementary Section S3 defines each measure in full, with its equation, what it rewards and what it cannot see; every one is quoted with the others rather than alone.

Statistical analysis

Held-out means and density strata are accompanied by 95% nonparametric bootstrap intervals (20,000 replicates; seed 20,260,810), resampled with stratification by density; the clean-minus-degraded analysis resamples paired scene differences within density. Per-coverage points in the comparison figures are plain means of ten scenes with no interval drawn and none claimed. The

Table 1. Held-out reconstruction scores overall and by target areal coverage. VI denotes variation of information, for which lower is better; higher is better for the other metrics.

Group	n	F_1	Precision	Recall	Recovery	ARI	VI split	VI merge	VI total
All	128	0.854	0.868	0.843	0.799	0.851	0.256	0.216	0.473
20%	32	0.911	0.928	0.896	0.918	0.907	0.137	0.093	0.230
30%	32	0.874	0.892	0.858	0.828	0.870	0.211	0.161	0.372
40%	32	0.858	0.864	0.853	0.796	0.855	0.242	0.224	0.466
60%	32	0.775	0.787	0.765	0.656	0.772	0.436	0.388	0.823

real SEM work is exploratory, 3 fields, and no inference is made from it; the inferential unit there is the field, and the 2566 reference crossings entering the crossing measures are a count pooled over six condition-specific evaluations of the three fields, not a sample size. Active-component ablations are one-at-a-time development experiments on all 84 development scenes and do not reuse the held-out set.

3 Results

3.1 Held-out synthetic reconstruction

PLECTA completed all 128 held-out scenes with 0 failures. Mean common-fragment pairwise F_1 was 0.854 (density-stratified 95% bootstrap CI [0.845, 0.864]); Table 1 gives precision, recall, reference-instance recovery, adjusted Rand index and both variation-of-information components, overall and by density. These scenes used unseen seeds from the same procedural generator as development data, so this is same-generator generalisation and not distribution transfer (Section 4.1).

Performance decreased with scene density: mean F_1 fell from 0.911 [0.888, 0.933] at 20% coverage to 0.775 [0.760, 0.790] at 60%, monotonically through the intermediate strata. The same decline appears on the paired clean-against-degraded set of Section 3.5, where the clean and the degraded curves separate at low coverage and meet at 50–60%: degradation costs least exactly where the score is lowest, because there the tangles rather than the gaps set the difficulty. Per-scene runtime had mean 2.54 s, median 1.12 s, 90th percentile 6.42 s, and maximum 34.06 s; these timings are hardware-specific.

3.2 Which density axis matters

Coverage is the fraction of the frame filled by the rendered bundles, so it rises both when filaments are added and when the same filaments are drawn thicker, and only the first changes the grouping problem. A factorial separates them: twelve geometries, three centreline length densities crossed with four seeds, were each rendered at bundle widths of 6, 11 and 16 px. Within one geometry the input axis mask and the overlap-aware centreline reference are byte-identical across the three widths, verified by SHA-256, while the union of the thick supports changes: areal coverage varies by up to a factor of 2.49. Any movement in the score across those three renderings would mean that the method or the metric reads the silhouette.

Across all 12 geometries the within-geometry F_1 range was 0.0000, that is, exactly zero in 12 of 12 (Table 2). Areal coverage therefore cannot causally affect the grouping score; PLECTA reads the axis mask, and the metric scores centreline fragments, so neither sees the width the bundles were drawn at. Centreline length density does move it, mean F_1 falling from 0.920 to 0.836 across the three densities of Table 2. The coverage bands of the three length densities also overlap, so coverage does not identify a difficulty level on its own. Density strata reported by target areal coverage should therefore be read as a proxy for the length density accompanying them under this generator.

Table 2. Bundle width against centreline length density, 36 development scenes. Within each row the three bundle widths share one byte-identical axis mask and one reference, so the spread in areal coverage carries no score difference. Length density is centreline pixels per image pixel.

Centreline length density	n	Areal coverage	Bundle width	F_1
0.0193	12	0.118–0.309	6, 11, 16 px	0.920
0.0372	12	0.218–0.505	6, 11, 16 px	0.878
0.0548	12	0.305–0.644	6, 11, 16 px	0.836

Table 3. All four methods on the degraded masks, identical input masks and one scorer. $n = 50$ scenes and 3638 reference crossings, except GraFT, over the 47 it completed. Variation of information is in bits and lower is better; higher is better for the rest. Crossing Fidelity is defined in Supplementary Section S3, and leaving every arm unpaired scores 0.088 on it here. Paired differences with their intervals are in the text.

Measure	PLECTA	Greedy	DNAi	GraFT
Common-fragment F_1	0.822	0.676	0.343	0.490
Precision	0.831	0.785	0.462	0.666
Recall	0.814	0.599	0.284	0.400
Reference-instance recovery	0.765	0.542	0.252	0.304
Adjusted Rand index	0.817	0.670	0.331	0.479
Variation of information $\downarrow$	0.573	1.036	2.107	1.558
Crossing Fidelity	0.809	0.729	0.581	0.683
Median s per scene, 20/60%	0.14–4.1	—	0.04–0.6	2.44–29.7

3.3 Four linkers on identical masks

One baseline and two published methods – DNAi [26] and GraFT [27], each run from its own released source – were scored on identical masks through the same conversion and metric code as PLECTA, over 50 scenes generated after every configuration was fixed. That path was gated on parity: no comparator number was believed until it reproduced PLECTA’s stored F_1 , adjusted Rand index and both variation-of-information components, because a faulty conversion yields plausible rather than obviously wrong scores. Each comparator’s parameter translation was selected on the same 84 development scenes PLECTA was configured on, and every adaptation was checked and found to favour the comparator; Supplementary Section S4 records each one and how each method fails.

Greedy continuation. The baseline skeletonises the same mask and greedily accepts the lowest-turning-angle pairing between branch ends inside distance and turning-angle gates fixed by a development sweep. It is ours – not a reimplement of any paper but of the pairing principle shared by SIFNE, GraFT’s path cover and Wegmayr et al. [23, 25, 27] – and is reported as a transparent floor, never as an external result. PLECTA leads it by +0.145 F_1 [+0.129, +0.162] paired over the 50 scenes, and by +0.128 [+0.113, +0.144] on clean axes, so the margin is not an artefact of mask degradation.

Exact against greedy matching, matcher only. That baseline replaces the whole linker; a second arm isolates the claim, keeping the graph, cost, admissibility gate and round schedule and replacing only the per-junction maximum-weight matching by greedy cheapest-first acceptance. Exactness was worth +0.080 F_1 [+0.070, +0.089] over 30 development scenes, on 29 of 30 of them, and +0.090 [+0.080, +0.100] on clean axes; the failure is division, not merging, with variation-of-information split rising 0.252 to 0.500 at unchanged merge, 30 % more instances emitted and crossings resolved exactly falling 0.861 to 0.668. The relative loss grows monotonically with degree, -21 % at degree three to -35 % at degree six, but in absolute terms peaks at degree four (-0.252), where most crossings lie; and the greedy arm runs at parameters




Figure 4. The three methods on the degraded input masks, by coverage stratum, under three of the measures of Section 2.7: **(a)** pairwise F_1 , which counts pairs of fragments; **(b)** detection F_1 at IoU 0.1 with the 2px placement tolerance, which counts objects; **(c)** Crossing Fidelity, which counts crossings, against the fidelity of leaving every arm unpaired (dashed). The panels share one 0–1 scale, with tick labels written once. The numbers inside (a) are each method’s variation-of-information total in bits at the two ends; that measure has no panel because PLECTA’s whole trajectory lies under 0.9 bits against DNAi’s 2.9 and a shared axis crushes it. Points are plain means of $n = 10$ scenes, no interval drawn and none claimed; the paired differences are in Table 3. GraFT exceeded its wall-clock budget on three of the densest scenes, the hardest of the set, so its rightmost point is a mean of seven and is optimistic in all three panels.




Figure 5. The same three methods, the same scenes, the same scorer and the same panels as Figure 4, from the generator’s undegraded axis masks instead of the degraded ones. Every method scores higher and the order does not change at any coverage under any of the three measures. One of GraFT’s densest scenes was censored here rather than three.

selected with the exact matcher in the loop, a bias in that matcher’s favour.

DNAi. Its segmentation network was bypassed and only its instance-forming stage used, so grouping is not confounded with segmentation quality, and it was run at a development-selected configuration rather than at its released defaults. That adaptation is worth a great deal to it: at those defaults it scores $F_1=0.195$ on these masks against 0.343 tuned. Tuned, the paired difference is $+0.478$ [$+0.456$, $+0.501$] over 50 scenes, and $+0.414$ [$+0.392$, $+0.436$] on clean axes, the gap widening monotonically with density (Figures 4 and 5).

DNAi’s margin is a limitation of domain and density, not a general ranking: it is built for DNA-fibre micrographs, which are sparse, and where the network is sparse and the axis clean it recovers 0.660 at 20% coverage. Restricting the score to the evidence a degree-limited, gap-blind method can reach separates capability from judgement, and on clean axes the difference falls from $+0.414$ to $+0.284$ [$+0.262$, $+0.307$], still favouring PLECTA in 49 of 50 scenes and at every

coverage level. Roughly 31 % of the margin is therefore coverage of cases the comparator cannot attempt, and 69 % the outcome of decisions both methods make.

GraFT. Its published entry point begins with a Frangi and Otsu front end for grayscale ridge images, which has nothing to do on a two-valued axis, so the mask was injected at its skeleton-to-graph stage, everything downstream being its own code unchanged. That substitution favours it: run end to end on the grayscale images it scores 0.100, because its segmentation stage recovers only a third to a half of the reference skeleton. Its minimum continuation angle was also re-tuned on development scenes. GraFT scored 0.490, a paired difference of +0.339 [+0.303, +0.378] favouring PLECTA in every one of the 47 scenes both methods completed, and +0.275 [+0.224, +0.329] on clean axes. 3 of the densest degraded scenes did not finish inside a 45-minute budget and are excluded from its means; they are the hardest scenes of the set, so that exclusion flatters GraFT. Neither released method establishes a general ranking, and Section 4.1 states what these comparisons leave untested.

A second generator. Every scene above comes from one generator, so the comparison was repeated on GraFT’s own, run unmodified at its published settings over 40 scenes in 8 density strata spanning 0.0016 to 0.0557 centreline pixels per image pixel; GraFT’s published comparison reaches 0.020, so its whole range is the lower half of this one. PLECTA leads at every density, from 0.970 against 0.787 and 0.659 at the sparsest to 0.800 against 0.219 and 0.149 at the densest; the margin widens with density and nothing crosses over, including inside GraFT’s published range. The field-standard measures order the three identically over the whole ladder (adjusted Rand index 0.912, 0.453 and 0.319; variation of information 0.262, 1.629 and 2.033 bits), and PLECTA’s errors lean mildly to fusion where both comparators divide far more than they fuse. GraFT was also run end to end with its own front end, which on this data is inside its design envelope rather than degenerate as it is on ours; that arm scored 0.609 and 0.237 at the two ends, below the injected arm almost everywhere, so injecting the mask continues to favour it. The held-out result of Section 3.1 is a same-generator measurement and remains one; this comparison is not.

Crossings resolved. Panel (c) of Figures 4 and 5 states the contribution directly, and pooled over the 3638 reference crossings of the degraded set PLECTA partitions 0.809 of them exactly against 0.729 for the greedy baseline, 0.683 for GraFT and 0.581 for DNAi; paired, PLECTA’s margins are +0.065 [+0.047, +0.083], +0.123 [+0.085, +0.165] and +0.185 [+0.158, +0.211]. Clean axes order them the same way, PLECTA at 0.827. Leaving every arm unpaired resolves 0.088 of these crossings, which is the all-unpaired baseline the rate is read against; it is a deterministic do-nothing policy rather than a chance level. It is much less discriminative than the pairwise score, spanning 0.259 between best and worst method at 60 % coverage where pairwise F_1 spans 0.536, so it is used to say what was solved and not to rank. Everyone is poor at the busiest nodes: above degree four PLECTA resolves 0.414 and no comparator does better, while DNAi collapses to 0.016, which is what its degree-four cap predicts.

Per-scene cost was measured warm, single-threaded and pinned to one core, timing each method’s own computation only (Table 3; 3 GraFT scenes exceeded 600 s and are censored). These are independently optimised codebases, so absolute times measure implementation as much as algorithm; the transferable quantity is scaling, and the log-log exponents against foreground pixels are 2.11, 2.65 and at least 2.01. PLECTA has the largest uncensored exponent; GraFT’s is a lower bound, because the scenes its fit lost to censoring are the slowest ones, and only GraFT is heavy-tailed.

3.4 Crossing representation, and where it over-marks

Resolving a crossing and representing one are different things, and only the first is what the scores above measure. Which arm continues into which is the difficult decision, and every score in this paper rewards getting it right. Whether the emitted raster *represents* a crossing pixel as belonging to both filaments is a property of the output format, and the metric cannot see




Figure 6. Reconstruction at three coverage levels, each the scene closest to the median F_1 of its stratum, shown as a horizontal window of the 512px scene rather than the whole of it. Columns: the degraded input mask, the overlap-aware reference, and PLECTA’s instances drawn at their fitted widths. What is evaluated in this paper is the grouping of centreline pixels the method produces from the mask alone, and the F_1 beside each row is that grouping’s score; the width rendering shown here is a downstream step that redraws those instances and cannot move a pixel from one to another. Instance colours are arbitrary and not matched between columns.

it: flattening the emitted layers to a single label image leaves PLECTA’s F_1 unchanged to four decimal places (0.0000).

The representation is instead verifiable directly: on the comparison set, reference instances share 11 % of their pixels with another instance at 20% coverage, rising to 39 % at 60%; PLECTA emits 5–20 %, and GraFT, whose output is close to a hard partition, about 2 %. Which pixels those are is less flattering. PLECTA recovers 0.463 of the pixels the reference gives to more than one instance, less than half, though the largest such recall here, against 0.144 for GraFT and 0.128 for DNAi, while the greedy baseline, a hard partition, scores zero by construction. It does so at a precision of 0.134, marking 3.6–4.0 times more shared pixels than the reference contains, because a chain is painted with every pixel of every junction cluster it touches. That is a property of the rendering rather than of the grouping, and one the F_1 neither rewards nor penalises.

Figure 6 shows what the score corresponds to scene by scene, taking at three coverage levels the scene whose F_1 is closest to the median, so the rows are typical rather than selected.

3.5 Response to the input mask

On the separate 50-scene controlled set, mean F_1 was 0.822 from degraded axes and 0.844 from the corresponding clean axes. The paired clean-minus-degraded change was 0.0224 (95% CI [0.0113, 0.0333]). Modelled gaps and raster irregularities therefore caused a small resolved loss under this generator.

The thickness of the axis mask costs more, and it matters because it is the upstream segmenter,

Table 4. One-at-a-time component ablations on the 84-scene development set. Negative ΔF_1 values denote losses relative to the evaluated configuration.

Development variant	F_1	ΔF_1	ARI
Fixed configuration	0.867	+0.000	0.864
No gap bridging	0.728	-0.139	0.722
Single round	0.847	-0.020	0.844
No chain-extended frames	0.846	-0.021	0.842
No round schedule	0.863	-0.004	0.860
No curvature term	0.863	-0.005	0.859
No short-connector node merging	0.845	-0.022	0.841
No junction-cluster consolidation	0.804	-0.063	0.799
No junction chord-turn term	0.681	-0.186	0.675

not the method, that fixes it: over 20 development scenes a nominal 3px axis mask scored -0.047 F_1 [-0.073, -0.023] against the 1px mask of the same scenes. Each width variant is scored against common fragments rebuilt from its own input mask, so the question asked is how well the method groups the skeleton it is handed, and the decline with width is a thicker mask skeletonising into a messier graph rather than a change in the grouping problem. This is the one density axis that does move the score. The rendered silhouette does not, causally and scene by scene (Section 3.2), so a mask should be reported by the geometry of its axis and not by how much of the frame its bundles fill.

3.6 Development ablations and downstream rendering

One-at-a-time ablations were evaluated only on the 84 development scenes (Table 4). The three largest losses relative to the evaluated configuration were, in order, the junction chord-turn term (-0.1858), gap matching (-0.1394) and connector-based junction-cluster consolidation (-0.0634); every other component, chain-extended frames and the round schedule included, mattered less. These are development values, not held-out effect estimates.

The core output is a grouping, not a set of connected curves: accepted gap links establish identity but are not painted, so on the development set 0.653 of core instances arrive as more than one disjoint pixel run and 0.663 contain a branch point. Any downstream step that assumes one connected curve per instance must account for this. The optional ribbon rendering of Section 2.5, at its defaults, scored development $F_1=0.860$, a change of -0.0068, and reduced those fractions to 0.000 and 0.005, so rendering buys connectivity at a small cost in the identity score without changing chain membership. Enabling the separate 0.6 silhouette-absorption rule reduced F_1 further to 0.857, and because absorption can merge identities it remained disabled. None of these topology measurements validates apparent width.

3.7 Proof of concept on real SEM

The claim of this section is one sentence: PLECTA works on real imaging when a second technique supplies the mask, and Pipeline A is such a technique. Nothing here is a measurement of Pipeline A, which carries its own evaluation [32].

Pipeline A runs procedural mask $\rightarrow$ CycleGAN $\rightarrow$ paired synthetic corpus $\rightarrow$ nnU-Net $\rightarrow$ binary filament-axis mask: a generator draws a mask, a CycleGAN translates it into an SEM-like image so that image and mask are paired without an annotator, an nnU-Net trained on those pairs segments a real micrograph, and its output is a binary axis mask [31, 32]. Pipeline B is PLECTA, and it is the only stage evaluated here. 3 manually annotated CNT-sheet fields separate the two: PLECTA is run once on the skeleton of the manual annotation and once on the nnU-Net axis, with the annotation as the reference in both cases, so the two conditions




Figure 7. The two input axes of field B58_110, drawn over the micrograph they were read from: (a) the skeleton of the manual annotation and (b) the output of Pipeline A. These are the binary masks themselves, before any reconstruction, so there is nothing to tell apart within a panel and one colour is sufficient; the comparison the figure makes is between them. Everything separating the two conditions of Table 5 is visible here: the same filaments, read a second time by a network rather than by a person, arrive thicker, broken into more pieces, and displaced by a pixel or two almost everywhere. The axis is dilated by one pixel for legibility at this size, which changes no geometry. B58_300 behaves the same way, as its mask-quality triple in the text records.

Table 5. Reconstruction from the manual-derived and the nnU-Net axis in three CNT SEM fields, under the measures of Section 2.7. *Held out* is whether the nnU-Net run that produced the axis excluded that field from training; B58_100 could not be held out by any available run, so its U-Net row is an upper bound. *CF* is Crossing Fidelity and *Unpaired* the fidelity of the all-unpaired baseline, the do-nothing policy that leaves every arm unpaired on that field. *Cross.* is the number of reference crossings scored. Detection F_1 is at IoU 0.1 with the 2px placement tolerance. Exploratory, 3 fields, no inference.

Field	Input axis	Held	Grouping			Objects and crossings			
		out	F_1	Join F_1	Decisions	Det. F_1	CF	Unpaired	Cross.
B58_100	Manual-derived	–	0.887	0.929	13 234	0.941	0.922	0.013	371
B58_100	U-Net	no	0.457	0.559	17 282	0.633	0.617	0.102	384
B58_110	Manual-derived	–	0.837	0.911	24 870	0.930	0.920	0.023	599
B58_110	U-Net	yes	0.452	0.593	16 134	0.602	0.614	0.048	290
B58_300	Manual-derived	–	0.947	0.959	28 224	0.959	0.956	0.011	611
B58_300	U-Net	yes	0.466	0.615	16 729	0.577	0.617	0.048	311

differ only in the mask. The first is an oracle-like input condition rather than an end-to-end result, asking what the grouping can do when the axis evidence is curated; only the second says what an automatic mask allows. Figures 8 and 7 show both on one field, at width and as the centrelines the scores are computed on.

Which fields are held out, and what the exception costs. B58_110 and B58_300 are absent from the training set of the inference run used here, verified by image content rather than by filename at tile correlations of 0.10 and 0.15, so both are genuine held-out observations of mask transfer. B58_100’s image is present in every available training run at 0.9999 or above and cannot be made otherwise; its U-Net row is an upper bound and is marked as such in Table 5. What that costs is measured rather than left vague. Running 3 nnU-Net models over all three fields scores the same field under models that held it out and one that did not: on B58_110 the two clean models give 0.451 and 0.453 against 0.512, on B58_300 0.464 and 0.468 against 0.537. Training on a field is worth 0.060 to 0.071 pairwise F_1 , the trained-on model predicting about a third more foreground, which is the mechanism. On B58_100 every model trained on the field, the three agree, and there is no gap to see: that is the control.

What the three fields say. From the manual-derived axis PLECTA reached $F_1=0.887$,




Figure 8. Reconstructed bundles on the two SEM fields held out of the nnU-Net run that produced their axes, one field per row. **(a)** the manual annotation; **(b)** PLECTA from the manual-derived axis; **(c)** PLECTA from the nnU-Net axis of Figure 7b, the same method and the same parameters, differing only in the mask. All three are drawn at width so that they are the same kind of object: the annotation as the annotator painted it, and the two reconstructions through the optional image layer of Section 2.5, which measures a width per instance across the centreline and renders a ribbon. That layer runs after the grouping is fixed and cannot move a pixel from one instance to another, so no score in this paper depends on it. Instance colours correspond across the panels of a row: a predicted instance takes the colour of the annotated filament it was matched to, one matched to none is crimson, and adjacent filaments are given different colours, which carry identity within the row and nothing else. Each crop is the window whose local pairwise F_1 is jointly closest to that field’s whole-field value on both axes, so it is typical of the field rather than selected from it.

0.837 and 0.947, and resolved 0.922, 0.920 and 0.956 of the reference crossings exactly against an all-unpaired baseline of 0.011 to 0.023. From the nnU-Net axis the same predictions read 0.452 and 0.466 on the two held-out fields, with Crossing Fidelity 0.614 and 0.617 and join F_1 0.593 and 0.615. The two clean fields agree with each other to within 0.03 on every measure, which is the substantive gain over the single clean field the study previously had; 3 fields still cannot estimate a population, and the inferential unit is the field: the 2566 reference crossings the crossing measures rest on are pooled over six condition-specific evaluations of these three fields, a pooled count and not a sample size.

Most of the distance between the two conditions is the mask, and the method fails by conservatism rather than by fusing. The axis each field hands PLECTA is characterised, before any reconstruction is opened, by the tolerant completeness/correctness/quality triple: at 2px, quality is 0.504, 0.557 and 0.507, over completeness 0.671, 0.651 and 0.619. The tolerance is not a formality: quality falls to 0.299, 0.375 and 0.325 at 1px, because the annotation and the nnU-Net output are two independent readings of the same filaments and disagree in placement almost everywhere. That displacement is also why detection F_1 carries a placement tolerance, and the control that makes it a correction rather than an inflation is here: on the manual-derived axes, where the input mask is a skeleton of the annotation itself, the tolerance moves detection F_1 by 0.000, against 0.179 rising to 0.602 on B58_110’s nnU-Net axis. Because the pairwise score is quadratic in fragments, a filament the nnU-Net mask breaks into several pieces carries every one of its pairs; join F_1 , which counts one decision per pair of fragment ends the mask brings within reach, reports 24870 decisions on B58_110’s manual axis against 16134 on its nnU-Net one. It narrows the gap without closing it, and its decomposition says what the rest is: a join precision of 0.681 against a recall of 0.525 is a method declining joins rather than fusing

Table 6. The depth stage on 12 depth-resolved synthetic scenes. Chance is 0.5 for every accuracy shown. Exploratory, six scenes per coverage, plain means, no inference.

Areal coverage	Given 2-D geometry		Reconstructed 2-D	
	20 %	40 %	20 %	40 %
Projected crossings per scene	28	150	28	150
Crossings recovered	1.000	1.000	0.857	0.746
Order accuracy, all crossings	0.831	0.808	0.725	0.588
Order accuracy, decided only	0.863	0.879	0.908	0.861
Abstained	0.038	0.082	0.069	0.087
Depth order, crossing pairs	0.858	0.864	0.888	0.819
Depth order, all pairs	0.640	0.710	0.641	0.696
Instance pairs sharing a component	0.392	0.852	0.340	0.751
Layer agreement	0.663	0.395	0.671	0.361

filaments, which is what panel (c) of Figure 8 shows.

3.8 Optional depth decomposition

The stages above reconstruct instances in projection. A crossing pixel is assigned to both filaments that own it, but nothing so far says which of the two passes in front. The optional depth stage answers that from the greyscale alone – brightness and edge sharpness across the crossing – then imposes one globally consistent vertical order and places each instance at a height at which nothing interpenetrates. It changes no 2-D instance.

It is evaluated on 12 depth-resolved synthetic scenes at two areal coverages, under the measures of Supplementary Section S5, in two input conditions. Given the reference 2-D geometry, the depth decision is measured alone; given only the binary axis mask, the same stage runs on PLECTA’s own reconstruction and carries its error. Table 6 reports both.

The decomposition is the result. Order accuracy over decided crossings spans only 0.048 (0.861–0.908) across both conditions and both densities, while accuracy over all crossings spans 0.243 (0.588–0.831) – and it falls by the amount the crossing recovery rate falls, 1.000 to 0.746. Once a crossing is found, the decision is made at close to the same accuracy whether the geometry was given or reconstructed, and whether the scene carries the sparser crossing load or roughly five times it. What limits the depth result is the 2-D stage not recovering every crossing, not the ordering of the crossings it does recover.

The global stack is weaker than the local decision, and the gap is a property of the evidence rather than of the decision. Depth order is recovered at 0.888 over pairs that cross but only 0.641 over all pairs, because depth propagates solely along crossings: two instances in different components of the crossing graph are related by no evidence whatever, and their order is a stated tie-break. That is measurable, not merely arguable. Only 0.340 to 0.852 of instance pairs lie within one component, and if the pairs that do are ordered at the crossing-pair rate while the rest are a coin flip, the all-pairs figure follows arithmetically: at 20 % coverage that predicts 0.632 against 0.641 measured. The account is not complete at the denser end, where the measured order falls below the prediction by up to 0.100, consistent with error accumulating along the long transitive chains that large components create. Depth claims are therefore scoped to the connected component, and a globally ordered film would need evidence that does not travel through crossings, such as defocus, or a second view, rather than a smoothness assumption, which these scenes would not support: their heights are drawn independently, so proximity carries no depth information by construction.

Layer agreement follows the same pattern, falling from 0.671 to 0.361 as coverage doubles




Figure 9. The depth stage on one synthetic scene, reconstructed from the binary axis mask. The image the scene was rendered from, that scene re-rendered from the reconstruction through a forward SEM model, and the film the reconstruction implies. The re-rendered panel is a consistency check on depths and radii, not a scored quantity: the difference Table 6 reports between input conditions is a difference in which crossings were *found*, which a rendered image cannot show. The table carries that measurement.




Figure 10. The same comparison on a real CNT SEM field, reconstructed from the manual annotation’s axis. **Nothing here is scored.** No depth ground truth exists for real SEM – the annotation records which pixels belong to which filament, not which filament passes in front – so this shows the stage runs end to end on real data and re-renders to something recognisable, not that the depths are right. The forward model’s appearance parameters are fitted to this micrograph, so the residual quoted beneath the panel is in-sample and is a sanity check rather than a measure of reconstruction quality.

while the recovered layer count stays close to the reference: the ordering is locally right, so the number of layers is about right, while small order errors move instances across layer boundaries.

4 Discussion

PLECTA addresses a narrow problem: assigning fragments of a binary filament-axis mask to overlapping projected instances. Explicit geometry recovers substantial identity there with no instance-labelled training data and no image intensity, and what recovers it is the combination of local and accumulated evidence, neither of which suffices alone: a junction is resolved by mutual compatibility rather than by the cheapest available continuation, and the evidence for that judgement is re-read along the chains as they assemble. That evidence is directional throughout, which fixes the condition the image must meet: a tangent measured on one side of a crossing

has to predict the arm leaving the other, which bending stiffness produces most directly and tension, confinement or bundling can produce as well. The ablations are consistent with that reading, the two largest losses being the junction chord-turn term and gap matching, both purely geometric. Where a filament is flexible on the scale of a junction the premise fails rather than degrades: there is then nothing local to match, and the correct response is to acquire evidence the projection has lost rather than to weaken the gates. The ablations locate that effect without proving each term independently necessary: the two large losses correspond to the two kinds of discontinuity, the junction and the missing mask segment, while accumulated context is worth distinctly less than either.

Difficulty follows the density of centrelines rather than the amount of ink, so what costs the score is the number of locally plausible alternatives at a crossing. The rest of the sensitivity belongs to the upstream segmenter, which fixes the thickness of the mask and therefore how clean a skeleton PLECTA is given. PLECTA can bridge a short omission, but it cannot infer a filament absent from the mask, nor reliably reject a false axis introduced upstream. Supplying identical masks isolates grouping in an algorithmic comparison; it does not establish end-to-end SEM performance, which is why the three-field case is a demonstration of the joined pipelines rather than a measurement of them. Ribbon rendering, being downstream of grouping, buys connectivity rather than identity, because the core paints arms and junction clusters but not the gap links that carry identity across a break; width and brightness can be reported for characterised instances, but their physical accuracy requires separate references.

4.1 Evidence boundary

The principal limitation is distributional, and it is narrower than a one-generator study. The held-out scenes use independent seeds but the same procedural generator family used for development, so their confidence intervals quantify sampling variation within that family and not transfer to another simulator, specimen, microscope, or segmentation pipeline. The *ranking*, however, does not rest on that family: GraFT’s own generator was run unmodified at its published settings over 40 scenes in 8 density strata, and its bundled frame separately, and the ordering held on both (Section 3.3). What is untested is transfer of the held-out *level* to a third generator, not of the comparison. The ablation and rendering studies are development-only and remain secondary to the fixed held-out result.

Within that family, one generator setting is worth singling out because the method’s grouping core is an angle-continuation argument and a reader may reasonably suspect the result of resting on it. The generator draws each centreline as a correlated-curvature random walk whose bending stiffness is one parameter, set to 0.035 rad per step because CNT bundles are stiff. Sweeping it from 50 % to 150 % of that value, which moves the median radius of curvature from 975 to 308 px in a 512 px frame, changes nothing that can be distinguished from noise: over 5 settings at two coverage targets and both mask conditions, 10 scenes each, the largest spread of level means is 0.048 against a smallest scene-to-scene standard deviation of 0.050. The headline numbers do not depend on this parameter across that band. The wavelength over which a bend is spread was swept separately, from 10 to 40 px about a default of 20, with the same result in 3 of 4 series; the fourth exceeds its own noise only marginally (0.076 against 0.052) and is non-monotonic with its minimum at the default, so it is read as noise. This is a robustness statement and not a scope bound: it locates no breaking point, and the sweep stays well inside the stiff regime the paper claims, reaching a 95th-percentile local bend of only 9.2° over 20 px.

The real SEM study is deliberately a case study with 3 fields, not a population validation. Manual-derived masks are an oracle-like input condition: they test grouping when the axis evidence is curated, not automatic SEM-to-instance performance. B58_100 is training-contaminated for the U-Net condition and cannot be made otherwise, so its row there is an upper bound; B58_110 and B58_300 are held out, and two fields that agree closely are a stronger statement than the one the study previously had, but they still cannot estimate generalisation. The two

mask sources also generate different fragment universes, so their F_1 values describe performance under each input condition rather than comparing identical objects. More independent fields, sealed before configuration, are still required.

Multi-reader annotation is no longer outstanding, and the result reframes the limitation above rather than softening it. A second reader independently annotated 3 of these fields. Judged against the reader whose skeleton it was given, PLECTA scores 0.786–0.947; judged against the other reader, 0.370–0.510. That collapse is not the method degrading. The two readers agree with each other at 0.361–0.535 by the same measure with no method in the loop, so the cross-reader scores sit at roughly the ceiling human agreement allows. The disagreement is systematic and not noise: the second reader exports 1.85–2.83 times as many objects, shorter, with 8–17 times as many components in the per-instance-skeleton union, and the directional variation of information is split-dominant one way and merge-dominant the other. That is two different answers to what an instance *is* – a whole physical filament, or a locally visible segment – not one reader being wrong. For the same reason the automatic-mask condition, at 0.383–0.466 against either reader, cannot be separated from reader disagreement on this evidence.

This is also the strongest available argument for evaluating on a generator at all. Instance identity through an occlusion is precisely the quantity manual annotation cannot fix: on real fields the reference is a reading, and two competent readings of the same micrograph differ by more than most of the method gaps this paper reports. A procedural scene is the only condition in which the answer is defined rather than adjudicated, because identity is assigned before rendering rather than recovered after it. The synthetic evaluation is therefore not a convenience standing in for real ground truth; it is the only place an unambiguous instance definition exists, and the real fields are correctly read as a case study of behaviour under a stated annotation protocol.

One caveat from the comparisons must be separated from the claim it seems to threaten. The comparators lose on how a crossing is *resolved*, which the metric scores and which is the difficult decision, not on how it is *represented*, which the metric cannot see: GraFT’s near-hard partition is a real limitation that costs it nothing here. Both properties are worth having, and Section 3.4 measures the representation directly, including where PLECTA over-marks it.

Two further limits concern the output itself. PLECTA recovers projected two-dimensional identity: shared pixels encode overlap but not which filament passes above another, so no three-dimensional topology follows. And the grouping metric evaluates ownership of centreline fragments, which is not a validation of rendered width; quantitative diameter claims require matched physical measurements across the relevant scale, contrast, and specimen ranges. Pairwise fragment scores also emphasise instances with many fragments, which is why recovery and the adjusted Rand index are reported alongside.

The comparisons themselves are bounded. The greedy rule is our reimplement of a principle rather than of any particular paper, and DNAi was run outside the domain it was built for, part of its cost function reading a colour channel our masks do not have. What the results support is narrow and mechanical: pairing endpoints greedily, or committing to a pairing with no option to decline, loses identity in dense crossings. They do not support a claim that PLECTA is generally superior; that would need several released implementations run in their own regimes by their own authors. Broader validation should extend to independent real masks and other filament domains, with input fragments and scoring rules held fixed.

5 Conclusions

PLECTA reconstructs overlapping filament instances from a binary axis mask by alternating exact junction matching with gated gap matching as local arm geometry expands into chain-level evidence. Its evidence is directional throughout, so it requires tangent and curvature to persist across a crossing at the scale of its fitting support, a condition bending stiffness contributes

to and rod-like filaments commonly meet. That is the regime of nanotube bundles and of the cytoskeletal and synthetic-fibre networks comparable methods target, and not that of freely flexible chains, where the projection retains nothing local to match. The fixed method achieved common-fragment $F_1 = 0.854$ on 128 independently seeded held-out scenes from the same procedural generator, with performance declining as network density increased, and partitioned 0.809 of the comparison set’s reference crossings exactly as the reference does. On a generator that is not ours, being GraFT’s own, run unmodified at its published settings over 40 scenes in 8 density strata, it reached $F_1 = 0.916$ against 0.475 and 0.349 for the two released comparators, so the ordering does not rest on the generator that produced the held-out set. On identical masks it gained +0.145 F_1 over a greedy minimum-turning-angle continuation rule, the rule shared by several published linkers, and +0.478 and +0.339 over two released implementations run from their own source; these support solving each junction exactly rather than pairing endpoints in priority order, but the released comparisons are bounded by domain and density and do not establish a general ranking. Development ablations identify junction chord-turn evidence, gap bridging, and junction-topology consolidation as the principal contributors, and optional image-informed width reconstruction improves geometric presentation while remaining separate from the grouping claim.

The proof of concept on carbon-nanotube sheets joins the two pipelines: an upstream synthetic-data-driven segmenter prepares the mask, and PLECTA resolves it into instances, demonstrated on 3 annotated SEM fields under manual-derived and nnU-Net masks, two of them held out from that segmenter, without establishing population generalisation or width accuracy. The method therefore offers an inspectable route to projected filament identity that draws on no random number generator and, in its grouping core, no training data; the demonstrated SEM route is not training-free, since the mask it consumes comes from a trained CycleGAN and nnU-Net upstream. Broader claims require independent real fields, validated physical widths, further released implementations run in a matched comparison, and an external-generator benchmark with sealed seeds.

CRediT author contribution statement

[To do: Fill in each author’s CRediT roles before submission, checked against actual contributions; keep the author names and order as listed above.]

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data and code availability

The PLECTA source, complete fixed configuration, compact seed manifest, and aggregate machine-readable evidence are available at <https://github.com/Jorallan/PLECTA>. Generated scenes, detailed per-scene records, and development diagnostics are retained locally. **[To do: Create a tagged archival release and add its DOI before submission; do not infer or preassign a DOI.]**

Acknowledgements

The authors thank the maintainers of NumPy, SciPy, scikit-image, NetworkX, and Matplotlib.

References

- [1] Michael F. L. De Volder, Sameh H. Tawfick, Ray H. Baughman, and A. John Hart. Carbon nanotubes: Present and future commercial applications. *Science*, 339(6119):535–539, 2013. doi: 10.1126/science.1222453.
- [2] Jonathan N. Coleman, Umar Khan, Werner J. Blau, and Yurii K. Gun’ko. Small but strong: A review of the mechanical properties of carbon nanotube-polymer composites. *Carbon*, 44(9):1624–1652, 2006. doi: 10.1016/j.carbon.2006.02.038.
- [3] Alexey N. Volkov and Leonid V. Zhigilei. Structural stability of carbon nanotube films: The role of bending buckling. *ACS Nano*, 4(10):6187–6195, 2010. doi: 10.1021/nn1015902.
- [4] Igor Ostanin, Roberto Ballarini, David Potyondy, and Traian Dumitrică. A distinct element method for large scale simulations of carbon nanotube assemblies. *Journal of the Mechanics and Physics of Solids*, 61(3):762–782, 2013. doi: 10.1016/j.jmps.2012.10.016.
- [5] Bernard K. Wittmaack, Abu Horaira Banna, Alexey N. Volkov, and Leonid V. Zhigilei. Mesoscopic modeling of structural self-organization of carbon nanotubes into vertically aligned networks of nanotube bundles. *Carbon*, 130:69–86, 2018. doi: 10.1016/j.carbon.2017.12.078.
- [6] Liyu Dong, Haibin Hang, Jin Gyu Park, Washington Mio, and Richard Liang. Detecting carbon nanotube orientation with topological analysis of scanning electron micrographs. *Nanomaterials*, 12(8):1251, 2022. doi: 10.3390/nano12081251.
- [7] Tamjid Imtiaz, Jacques Doumani, Fuyang Tay, Natsumi Komatsu, Stephen Marcon, Motonori Nakamura, Saunab Ghosh, Andrey Baydin, Junichiro Kono, and Ahmed Zubair. Facile alignment estimation in carbon nanotube films using image processing. *Signal Processing*, 202:108751, 2023. doi: 10.1016/j.sigpro.2022.108751.
- [8] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-Net: Convolutional networks for biomedical image segmentation. In *Medical Image Computing and Computer-Assisted Intervention – MICCAI 2015*, volume 9351 of *Lecture Notes in Computer Science*, pages 234–241, 2015. doi: 10.1007/978-3-319-24574-4_28.
- [9] Fabian Isensee, Paul F. Jaeger, Simon A. A. Kohl, Jens Petersen, and Klaus H. Maier-Hein. nnU-Net: A self-configuring method for deep learning-based biomedical image segmentation. *Nature Methods*, 18:203–211, 2021. doi: 10.1038/s41592-020-01008-z.
- [10] Carsen Stringer, Tim Wang, Michalis Michaelos, and Marius Pachitariu. Cellpose: A generalist algorithm for cellular segmentation. *Nature Methods*, 18:100–106, 2021. doi: 10.1038/s41592-020-01018-x.
- [11] Uwe Schmidt, Martin Weigert, Coleman Broadus, and Gene Myers. Cell detection with star-convex polygons. In *Medical Image Computing and Computer Assisted Intervention – MICCAI 2018*, *Lecture Notes in Computer Science*, pages 265–273, 2018. doi: 10.1007/978-3-030-00934-2_30.
- [12] Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, Tete Xiao, Spencer Whitehead, Alexander C. Berg, Wan-Yen Lo, Piotr Dollár, and Ross Girshick. Segment anything. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 3992–4003, 2023. doi: 10.1109/ICCV51070.2023.00371.
- [13] Carsten Steger. An unbiased detector of curvilinear structures. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 20(2):113–125, 1998. doi: 10.1109/34.659930.
- [14] Alejandro F. Frangi, Wiro J. Niessen, Koen L. Vincken, and Max A. Viergever. Multiscale vessel enhancement filtering. In *Medical Image Computing and Computer-Assisted Intervention – MICCAI’98*, volume 1496 of *Lecture Notes in Computer Science*, pages 130–137, 1998. doi: 10.1007/BFb0056195.
- [15] Yoshinobu Sato, Shin Nakajima, Nobuyuki Shiraga, Hideki Atsumi, Shigeyuki Yoshida, Thomas Koller, Guido Gerig, and Ron Kikinis. Three-dimensional multi-scale line filter for segmentation and visualization of curvilinear structures in medical images. *Medical Image Analysis*, 2(2):143–168, 1998. doi: 10.1016/S1361-8415(98)80009-1.
- [16] Andrew M. Stein, David A. Vader, Louise M. Jawerth, David A. Weitz, and Leonard M. Sander. An algorithm for extracting the network geometry of three-dimensional collagen gels. *Journal of Microscopy*, 232(3):463–475, 2008. doi: 10.1111/j.1365-2818.2008.02141.x.
- [17] Jeremy S. Bredfeldt, Yuming Liu, Carolyn A. Pehlke, Matthew W. Conklin, Joseph M. Szulcowski, David R. Inman, Patricia J. Keely, Robert D. Nowak, Thomas R. Mackie, and Kevin W. Eliceiri. Computational segmentation of collagen fibers from second-harmonic generation images of breast cancer. *Journal of Biomedical Optics*, 19(1):016007, 2014. doi: 10.1117/1.JBO.19.1.016007.
- [18] Ting Xu, Dimitrios Vavylonis, Feng-Ching Tsai, Gijssje H. Koenderink, Wei Nie, Eddy Yusuf, I-Ju Lee, Jian-Qiu Wu, and Xiaolei Huang. SOAX: A software for quantification of 3D biopolymer networks. *Scientific Reports*, 5:9081, 2015. doi: 10.1038/srep09081.
- [19] Daniel A. D. Flormann, Moritz Schu, Emmanuel Terriac, Divyendu Thalla, Lucina Kainka, Marcus Koch, Annika K. B. Gad, and Franziska Lautenschläger. A novel universal algorithm for filament network tracing and cytoskeleton analysis. *The FASEB Journal*, 35(5):e21582, 2021. doi: 10.1096/fj.202100048R.
- [20] Nathan A. Hotaling, Kapil Bharti, Haydn Kriel, and Carl G. Simon. DiameterJ: A validated open source nanofiber diameter measurement tool. *Biomaterials*, 61:327–338, 2015. doi: 10.1016/j.biomaterials.2015.05.015.

- [21] Saurav Basu, Chi Liu, and Gustavo Kunde Rohde. Extraction of individual filaments from 2D confocal microscopy images of flat cells. *IEEE/ACM Transactions on Computational Biology and Bioinformatics*, 12(3):632–643, 2015. doi: 10.1109/TCBB.2014.2372783.
- [22] David Breuer and Zoran Nikoloski. DeFiNe: An optimisation-based method for robust disentangling of filamentous networks. *Scientific Reports*, 5:18267, 2015. doi: 10.1038/srep18267.
- [23] Zhen Zhang, Yukako Nishimura, and Pakorn Kanchanawong. Extracting microtubule networks from superresolution single-molecule localization microscopy data. *Molecular Biology of the Cell*, 28(2):333–345, 2017. doi: 10.1091/mbc.E16-06-0421.
- [24] Jaydeep De, Li Cheng, Xiaowei Zhang, Feng Lin, Huiqi Li, Kok Haur Ong, Weimiao Yu, Yuanhong Yu, and Sohail Ahmed. A graph-theoretical approach for tracing filamentary structures in neuronal and retinal images. *IEEE Transactions on Medical Imaging*, 35(1):257–272, 2016. doi: 10.1109/TMI.2015.2465962.
- [25] Viktor Wegmayr, Aytunc Sahin, Björn Sæmundsson, and Joachim M. Buhmann. Instance segmentation for the quantification of microplastic fiber images. In *2020 IEEE Winter Conference on Applications of Computer Vision (WACV)*, pages 2199–2206, 2020. doi: 10.1109/WACV45572.2020.9093352.
- [26] Clément Ployat, Yosra Mehrjoo, Renaud Duval, Marie Carole Boucher, Santiago Costantino, and Hugo Wurtele. DNAi: An open-source AI tool for unbiased DNA fiber analysis. *Nucleic Acids Research*, 54(7):gkag335, 2026. doi: 10.1093/nar/gkag335.
- [27] Isabella Østerlund, Arne Neumann, Zhiming Ma, Yansong Miao, Staffan Persson, and Zoran Nikoloski. GraFT: A robust network-based spatiotemporal analysis of filamentous structures. *Science Advances*, 12(13):eadz4132, 2026. doi: 10.1126/sciadv.adz4132.
- [28] Yi Liu, Abhishek Kolagunda, Wayne Treible, Alex Nedo, Jeffrey Caplan, and Chandra Kambhampettu. Intersection to overpass: Instance segmentation on filamentous structures with an orientation-aware neural network and terminus pairing algorithm. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 125–133, 2019. doi: 10.1109/CVPRW.2019.00021.
- [29] Max Frei and Frank Einar Kruis. FibeR-CNN: Expanding Mask R-CNN to improve image-based fiber analysis. *Powder Technology*, 377:974–991, 2021. doi: 10.1016/j.powtec.2020.08.034. Published online 2020.
- [30] Yi Liu, Su Peng, Jeffrey Caplan, and Chandra Kambhampettu. Pick and trace: Instance segmentation for filamentous objects with a recurrent neural network. In *Medical Image Computing and Computer Assisted Intervention – MICCAI 2023*, Lecture Notes in Computer Science, pages 635–645, 2023. doi: 10.1007/978-3-031-43993-3_61.
- [31] Bastian Rühle, Julian Frederic Krumrey, and Vasile-Dan Hodoroba. Workflow towards automated segmentation of agglomerated, non-spherical particles from electron microscopy images using artificial neural networks. *Scientific Reports*, 11(1):4942, 2021. doi: 10.1038/s41598-021-84287-6.
- [32] Oday Allan. CNT_SEM: CycleGAN synthesis and U-Net segmentation workflows for carbon-nanotube SEM data. GitHub repository, 2026. URL https://github.com/Jorallan/CNT_SEM.
- [33] Torben Peters, John Schumann, Kerstin Kämpf, and Asmus Meyer-Plath. AI-assisted recognition of fibre instances in scanning electron microscope images for aerosol characterization and exposure assessment. *Powder Technology*, 469:121909, 2026. doi: 10.1016/j.powtec.2025.121909.
- [34] Cian Gabbett, Luke Doolan, Kevin Synnatschke, Laura Gambini, Emmet Coleman, Adam G. Kelly, Shixin Liu, Eoin Caffrey, Jose Munuera, Catriona Murphy, Stefano Sanvito, Lewys Jones, and Jonathan N. Coleman. Quantitative analysis of printed nanostructured networks using high-resolution 3D FIB-SEM nanotomography. *Nature Communications*, 15:278, 2024. doi: 10.1038/s41467-023-44450-1.
- [35] Guillem Palou and Philippe Salembier. Monocular depth ordering using T-junctions and convexity occlusion cues. *IEEE Transactions on Image Processing*, 22(5):1926–1939, 2013. doi: 10.1109/TIP.2013.2240002.
- [36] Hyunmin Lee and Jaesik Park. Instance-wise occlusion and depth orders in natural scenes. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 21210–21221, 2022.
- [37] Menglei Chai, Lvdi Wang, Yanlin Weng, Yizhou Yu, Baining Guo, and Kun Zhou. Single-view hair modeling for portrait manipulation. *ACM Transactions on Graphics*, 31(4):116:1–116:8, 2012. doi: 10.1145/2185520.2185612.
- [38] Lawrence Hubert and Phipps Arabie. Comparing partitions. *Journal of Classification*, 2(1):193–218, 1985. doi: 10.1007/BF01908075.
- [39] Marina Meilă. Comparing clusterings—an information based distance. *Journal of Multivariate Analysis*, 98(5):873–895, 2007. doi: 10.1016/j.jmva.2006.11.013.
- [40] Todd A. Gillette, Kerry M. Brown, and Giorgio A. Ascoli. The DIADEM metric: Comparing multiple reconstructions of the same neuron. *Neuroinformatics*, 9(2-3):233–245, 2011. doi: 10.1007/s12021-011-9117-y.
- [41] Elizabeth P. Reilly, Jeffrey S. Garretson, William R. Gray Roncal, Dean M. Kleissas, Brock A. Wester, Mark A. Chevillet, and Matthew J. Roos. Neural reconstruction integrity: A metric for assessing the connectivity accuracy of reconstructed neural networks. *Frontiers in Neuroinformatics*, 12:74, 2018. doi: 10.3389/fninf.2018.00074.